

FIT9132 2017 S2

FIT9132 Introduction to Databases

FIT Database Teaching Team

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

FIT9132 2017 S2

FIT9132 Introduction to Databases

FIT Database Teaching Team

License: Copyright © Monash University. All Rights Reserved.

Generated by [Alexandria](https://www.alexandriarepository.org) (https://www.alexandriarepository.org) on August 18, 2017 at 4:53 pm AEST

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Contents

Title	i
Copyright	ii
1 Week 1: Introduction to SQL Developer	1
1.1 Connecting to Oracle database using SQL Developer	13
1.2 Using SQL Developer GUI to manage data	16
1.3 Data Anomalies	19
2 Week 2: The Relational Data Model FIT9132	21
3 Week 3: Conceptual Modelling	24
3.1 Using Tools to draw ERD	26
3.2 Building Conceptual Models	29
4 Week 4: Logical Modelling	31
4.1 Logical Modelling - Task B - Rental Model	45
4.2 SQL Developer Data Modeller Issues	46
5 Week 5: Normalisation	49
6 Week 6: SQL Data Definition Language (DDL)	53
6.1 INSERTing data into the database	57
7 Week 7: SQL Part I - SQL Basic	62
8 Week 8: Maintaining data	65
8.1 Transactions Management	67
9 Week 9: SQL Part II - SQL Intermediate	69
10 Week 10: SQL Part III - SQL Advanced	71
11 Week 11: PL/SQL Triggers	73

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

1

Week 1: Introduction to SQL Developer

The exercises in this first tutorial will introduce you to the SQL Developer software that you will use to connect to the Oracle database. The main aim of these exercises is to get you familiar with the SQL Developer interface/menu/options. At the end of this exercise, you should be able to:

1. make a connection and login to your account on the Monash Oracle database
2. change your Oracle password, and
3. disconnect from Monash Oracle database in SQL Developer.

There are a number of different ways to connect to an Oracle database. In this unit, you will use Oracle software called SQL Developer.

The SQL Developer software can be downloaded from your unit's Moodle site. Information on how to install and configure SQL Developer is available on the Moodle site under Software and Documentation section. Please also note that accessing the Oracle database from a machine located outside Monash University's network will require you to connect to the Monash VPN (Virtual Private Network) service. Information on how to install the Monash VPN client and then connect to the VPN service can be found at <http://intranet.monash.edu.au/esolutions/software/install-vpn.pdf>

Assignment Project Exam Help

Before looking at how we connect to the Oracle database and run SQL using SQL Developer, please take a few moments and configure SQL Developer so that the software will satisfy the needs of this unit. To do this follow the steps shown below. The steps are also available in the "SQL Developer Preference Settings" document (supplied on your unit Moodle site under the "Unit Information/Software Documentations" heading).

<https://powcoder.com>

Add WeChat powcoder

SQL Developer Preference Settings

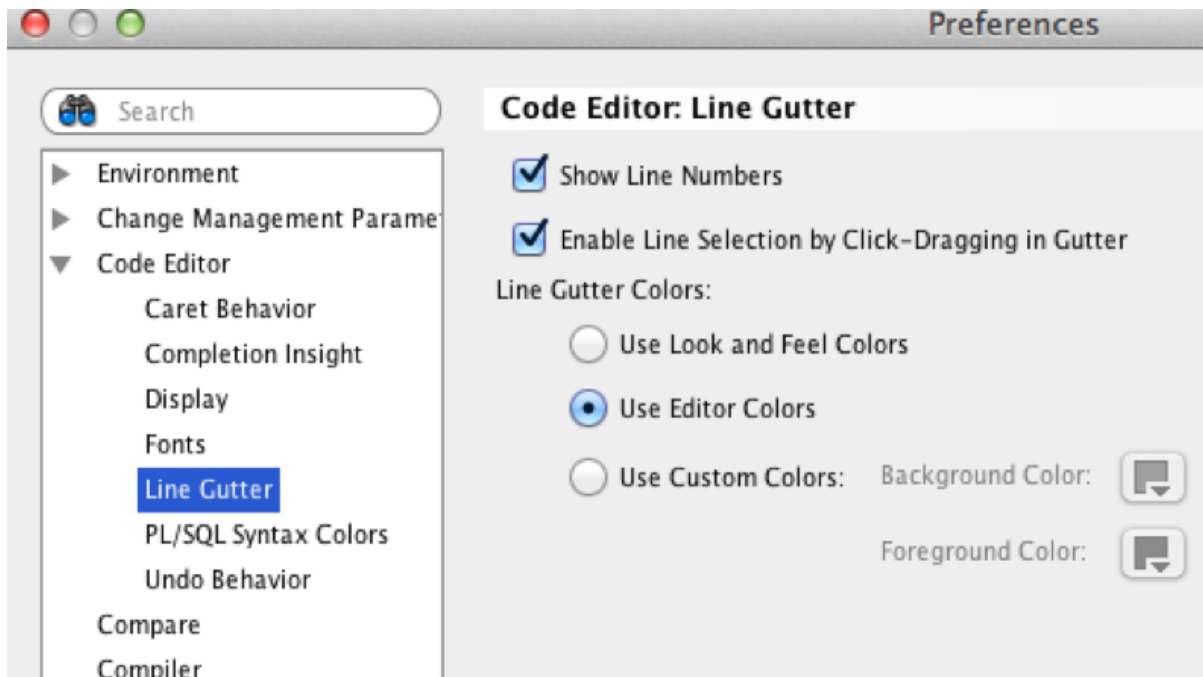
To help with your studies through the remainder of the semester you should configure SQL Developer to:

- display line numbers, and
- auto format SQL Code

This document also contains a number of other important configurations that you should implement before working extensively with SQL Developer. Under Windows the SQL Developer preferences are accessed from the Tools menu, under MacOS from the Oracle SQL Developer menu on the left.

Display Line Numbers

To do this select: Preferences - Code Editor - Line Gutter - Show Line Numbers (Check Show Line Numbers):



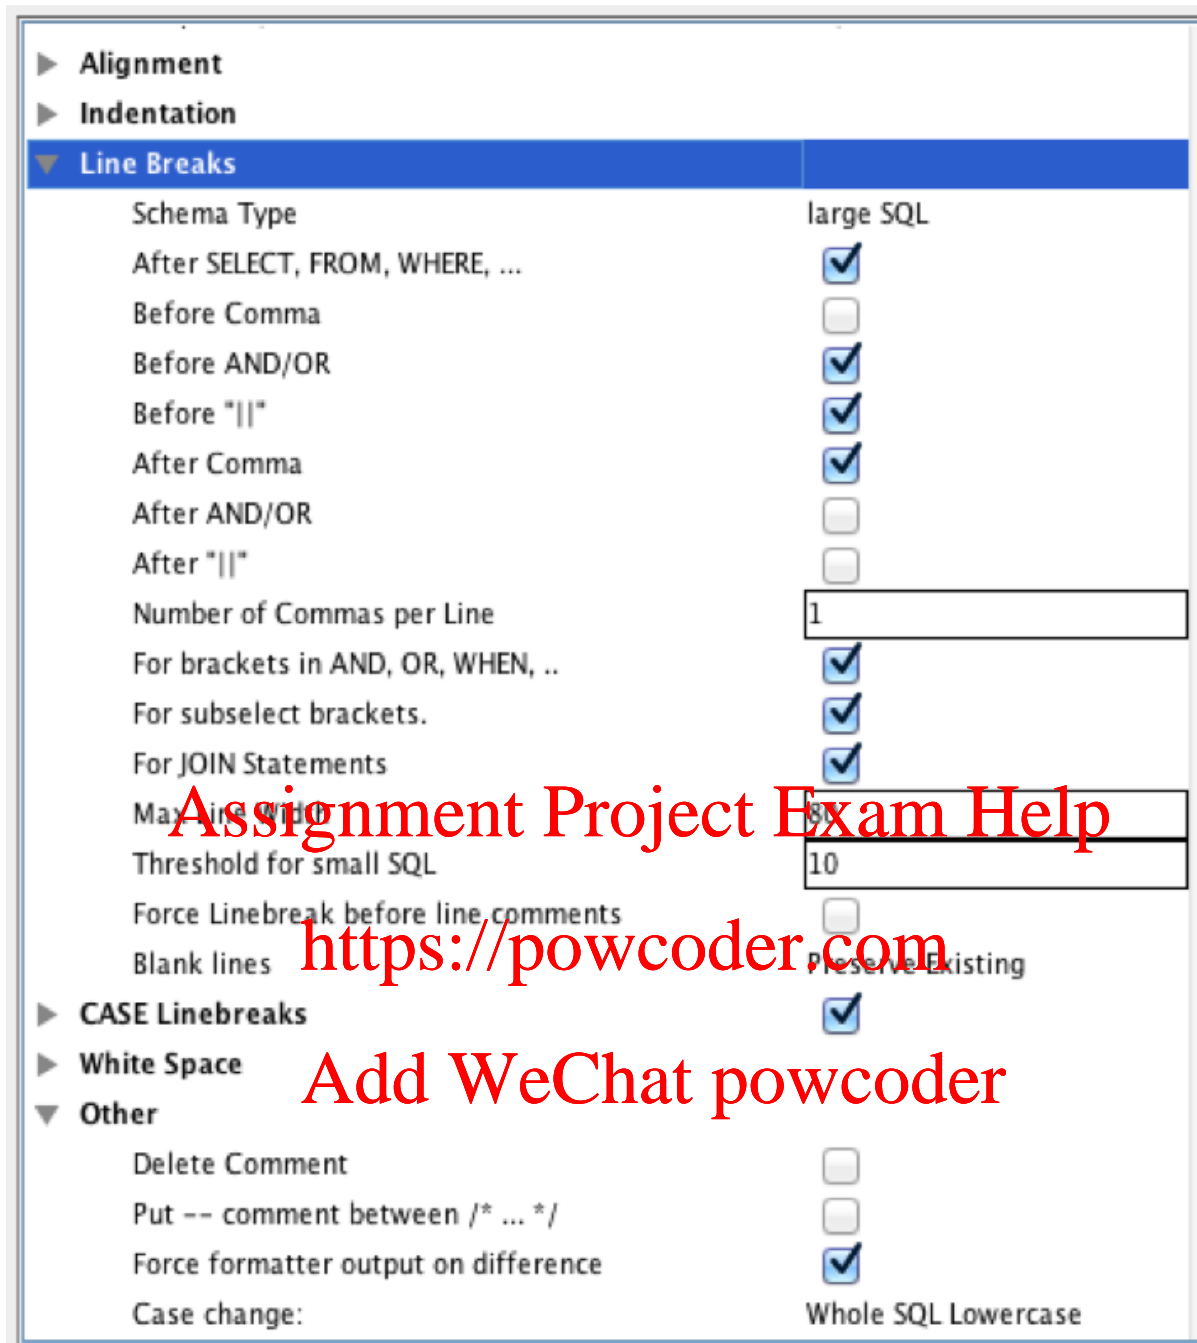
SQL Code Auto Format

We now wish to configure SQL Developer so that it will reformat any SQL we enter into a well set out "pretty" format. The aim will be to simply type SQL commands without worrying about layout and then use the SQL Developer assigned format key (see below) to automatically format the code.

<https://powcoder.com>

Version 4.1: Preferences - Database - SQL Formatter - Oracle Formatting - Select the SQL Profile from the drop down Profile menu and then select edit. Modify "Threshold for small SQL, under Line Breaks, to a small value such as 10:

Add WeChat powcoder

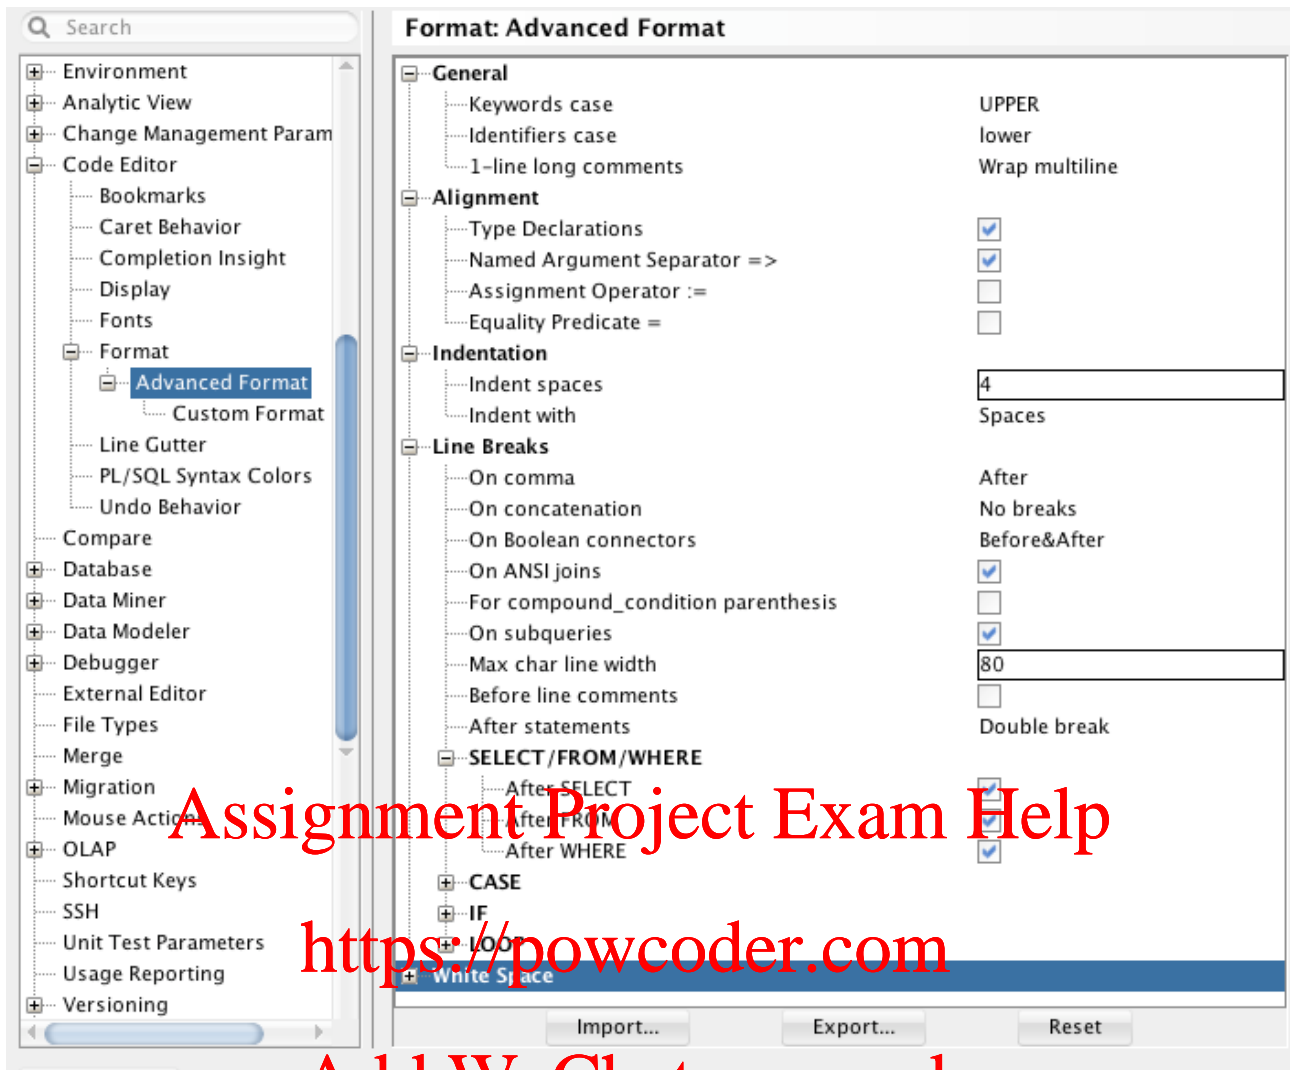


It is suggested that you also:

- set the "Max Line Width" to some reasonable value eg. 80 or 100,
- select (tick) the items selected above, and
- set "Blank lines" to "Preserve Existing"

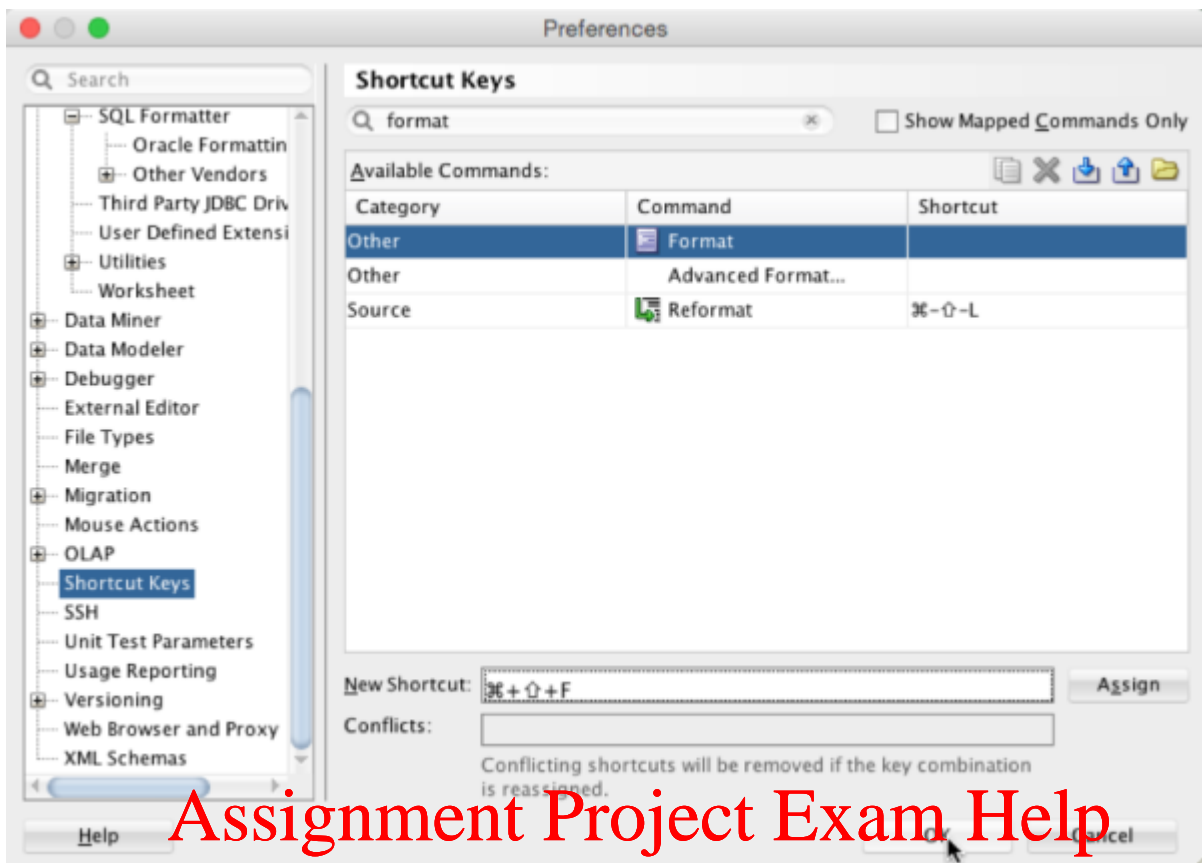
Depending on your personal preferences you might like to set "Case Change" under "Other" to "Whole SQL Lowercase" or leave it as the default of "Keywords Uppercase". Select OK

Version 4.2 & 17.2: Preferences - Code Editor - Format - Advanced Format.



For either version, you are free to use any settings which you are happy with; the settings shown are a starting point for your study.

To automate formatting we now need to assign a format key combination - select Preferences - Short Cut Keys and set a shortcut key for "Format". Your installation may already have a key set for format (eg. on Windows it is Ctrl+F7) - if you are happy with that key you do not need to change anything. Below shows setting a shortcut key (command+shift+F) for a mac:



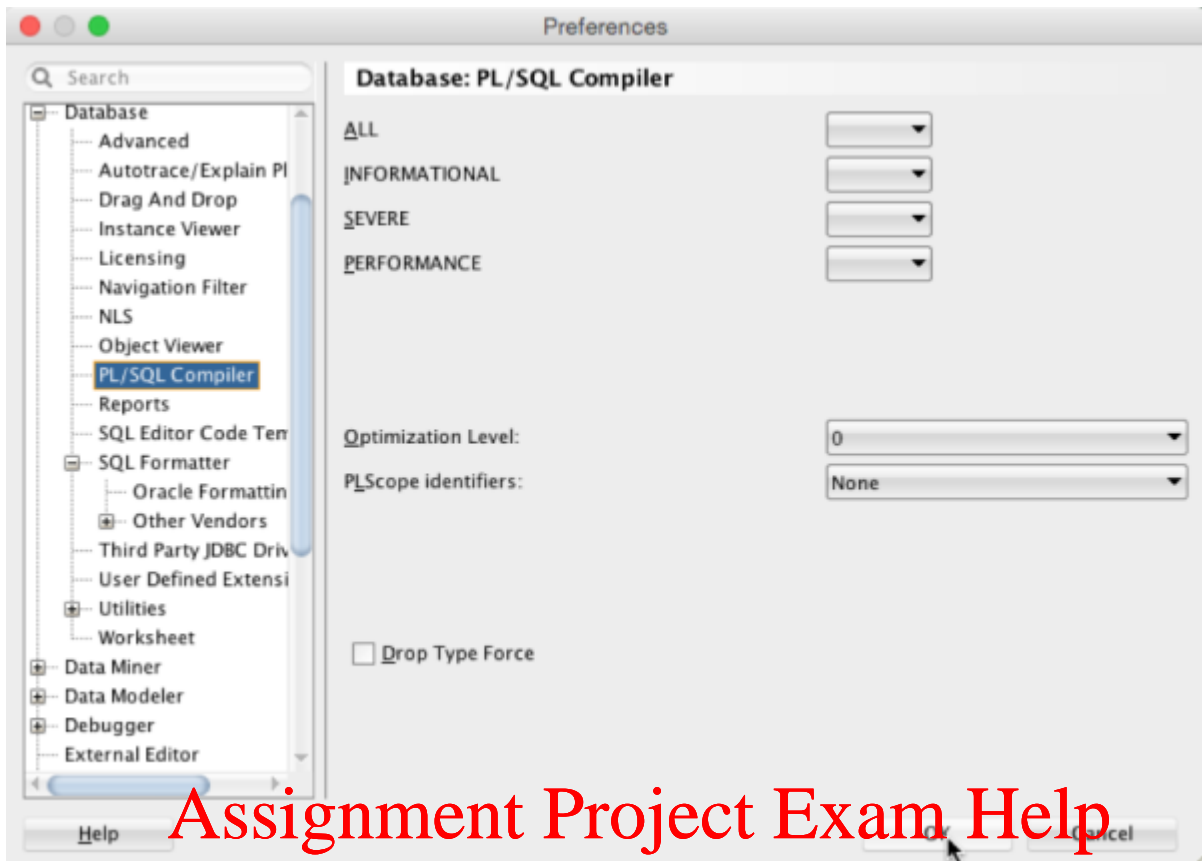
Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

PL/SQL Scope Identifiers

Please ensure you set the PL/Scope identifiers, under Database, to "None":



Assignment Project Exam Help

<https://powcoder.com>

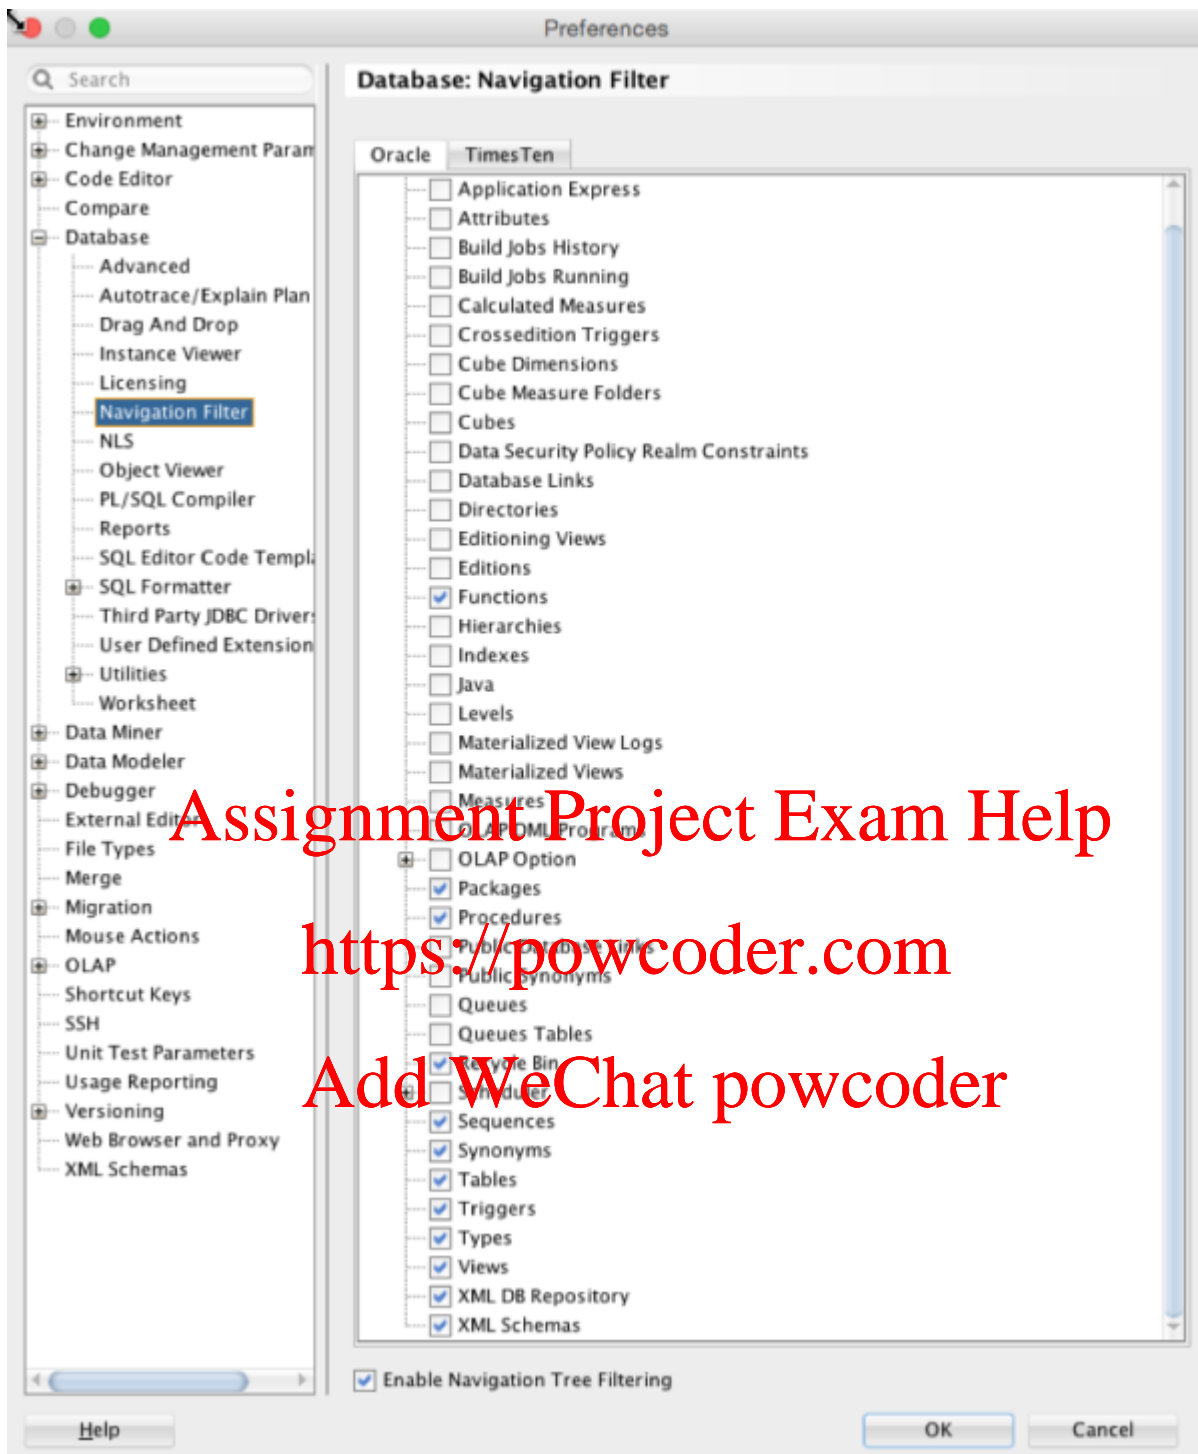
Navigation Filter (Optional)

SQL Developer displays a large, and often confusing array of items under its connection tree, these can be limited by applying a navigation filter (Database -> Navigation Filter).

Add WeChat powcoder

Select "Enable Navigation Tree Filtering"

Select to show nothing on the "Times Ten" tab. The suggested items for the Oracle tab are:



Copying your configuration to a new PC/Laptop

The configurations you have completed can be transferred to a different computer by manually copying the file `product-preferences.xml`. This process is also useful in the Monash on-campus labs if your PC has been configured to reset the configuration with each login - simply save the preferences file to your mapped desktop folder (see the section below 'Working in the On Campus Labs' on how to create this mapping).

Simply copy the file `product-preferences.xml` from the configured version to the new version where you want the same settings:

On windows the files is located in the folder:

C:\Users\yourusername\AppData\Roaming\SQL Developer\systemx.x.x.x\o.sqldeveloper.x.x.x.x\product-preferences.xml

for the current SQL Developer version and a user lsmi1 this would be:

C:\Users\lsmi1\AppData\Roaming\SQL Developer\system4.1.5.21.78\o.sqldeveloper.12.2.0.21.78\product-preferences.xml

If the AppData folder does not appear under C:\Users\yourusername you will need to [turn show hidden files and folders on](https://support.microsoft.com/en-au/help/14201/windows-show-hidden-files) (<https://support.microsoft.com/en-au/help/14201/windows-show-hidden-files>).

On the mac the files is located in the folder:

/Users/yourusername/.sqldeveloper/systemx.x.x.x.x/o.sqldeveloper.x.x.x.x.x/product-preferences.xml

for the current SQL Developer version and a user lsmi1 this would be:


/Users/lsmi1/.sqldeveloper/system4.1.5.21.78/o.sqldeveloper.12.2.0.21.78/product-preferences.xml

You are strongly recommended to make a copy of this file, once you are happy with your settings, and keep it in a safe place.

<https://powcoder.com>

To open the /Users/yourusername/.sqldeveloper folder on your mac, in Finder select from the Finder menu Go - Go to Folder and enter the path ~/.sqldeveloper :

Add WeChat powcoder



SQL Developer Language

If your laptop or computer is setup to use a non-English language (eg. Chinese), when SQL Developer installs it will make use of this system language as its default language.

For your Monash study, you **must modify** this so that the language being used (and displayed) by SQL Developer is English. To achieve this, edit the sqldeveloper.conf file - it's normal location is:

MS Windows:

C:\Program Files\sqldeveloper\sqldeveloper\bin\sqldeveloper.conf

OSX:

/Applications/SQLDeveloper.app/Contents/Resources/sqldeveloper/sqldeveloper/bin/sqldeveloper.conf

and add the line:

AddVMOption -Duser.language=en

Note that you **must carefully type (case is important) this line into the sqldeveloper.conf file yourself**, do not copy and paste it from this document. This entry must be on a line by itself and left aligned (add a new line at the end).


Also, note this must be done with a text editor. For MS Windows please do not use Windows notepad instead use something like [Notepad++](http://notepad-plus-plus.org/) (<http://notepad-plus-plus.org/>)

Working in the On-Campus Labs

The technique which the University uses in the on-campus Windows 10 labs to allow your files and desktop to be portable between different machines causes SQL Developer considerable problems in saving files.

To save files (.sql, .dmd, etc) from SQL Developer on the Monash Windows 10 on-campus labs we must use a mapped drive. If you have not as yet created such a drive, please use the following procedure.

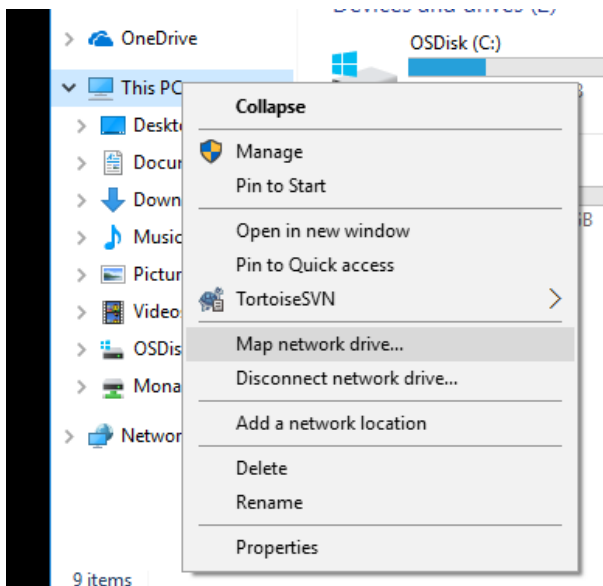
First, create a folder on your Desktop (say MyFolder, or you can use any folder you have previously created), navigate to this folder and then right click in the path bar of the Windows file manager on the folder, and select "Copy address as text".



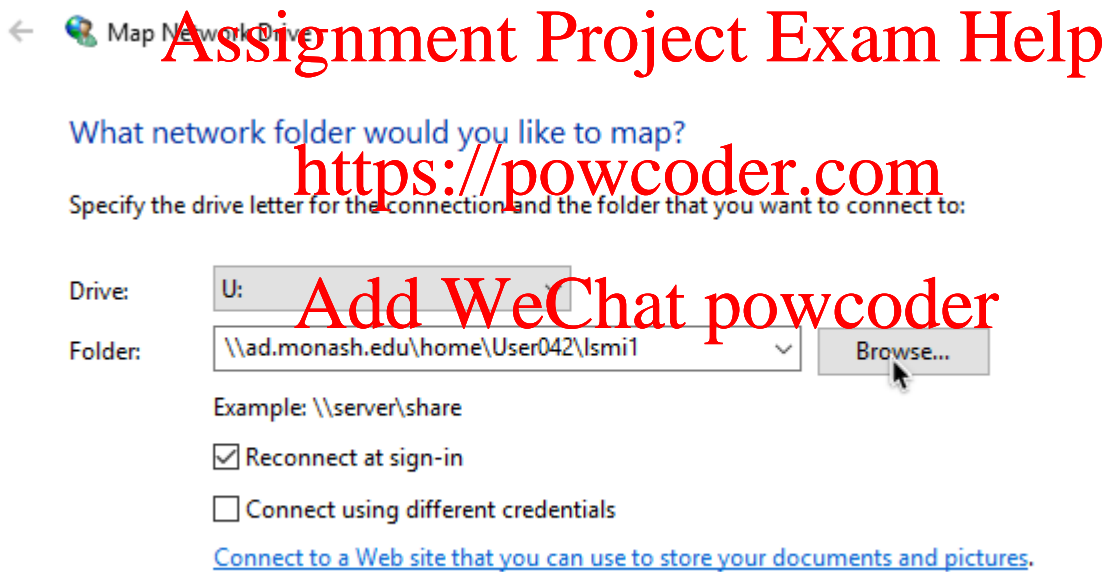
This will result in an address of the form:

\\ad.monash.edu\home\User042\lsmi1\Desktop\MyFolder


You should then use the first part of this path (exclude the ending \Desktop\MyFolder) to map a drive in Windows (use any drive letter you wish) however, we suggest U. First right click This PC and select "Map Network Drive":



Then enter the path that you copied above (without the Desktop eg. \\ad.monash.edu\home\User042\lsmi1) into the Folder entry:



Be sure to check "Reconnect at sign-in", then select Browse



and add a new Folder via "Make New Folder" under your authcate username. We suggest you call the folder "Units" - you can then create subfolders below this (after we have mapped the drive) for the various units you will study throughout your course. This setup only needs to be done once.

Assignment Project Exam Help

Finally, after creating the folder and selecting OK:

What network folder would you like to map?

<https://powcoder.com>

Specify the drive letter for the connection and the folder that you want to connect to:

Drive:

Folder:


Example: \\server\share

☒ Reconnect at sign-in

☐ Connect using different credentials

[Connect to a Web site that you can use to store your documents and pictures.](#)


Add WeChat powcoder



select Finish (double check "Reconnect at sign-in" is selected first). You now will have a drive U: below which you can create unit based subfolders. This approach has the advantage that the files you save on drive U: will be saved immediately to the drive. If you save anything to the Desktop these files are only transferred when you log out and can substantial slow down your login and logout times.

Please note that students currently have a space limit of 5Gb on this sever.

You will then be able to load and save files in the on-campus labs from your mapped drive eg. U: and access this drive from *any* University Windows Lab PC that you login to. To reach the drive you may need to pull down the SQL Developer folders list:



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

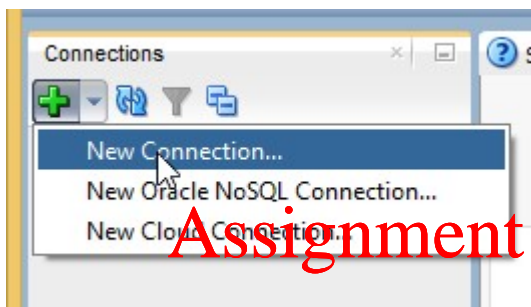
1.1

Connecting to Oracle database using SQL Developer

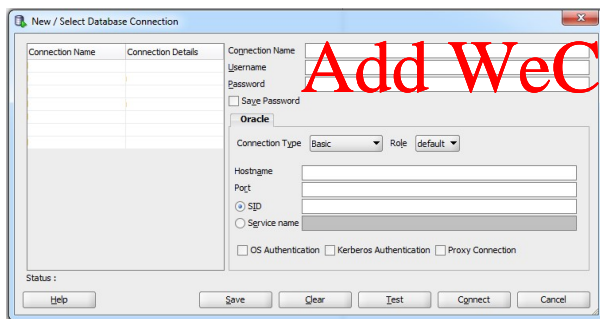
In the next few sections, you will learn how to use your newly configured SQL Developer software to access an Oracle database.

1. Adding a new connection

After running SQL Developer, right click on the Connections icon in the left panel, as shown below and select the New Connection option.

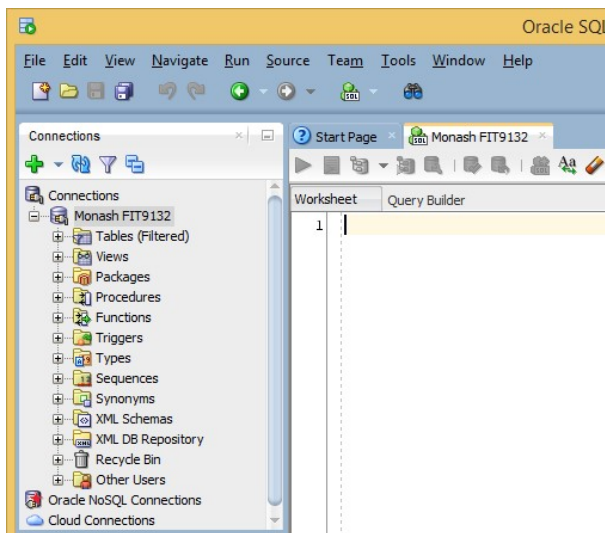


Now you will see the New / Select Database Connection window, as shown below.



The connection details that you need to connect to the Oracle database through SQL Developer will be provided by your tutor/lecturer. As well as setting up this connection, it is very important that you configure SQL Developer correctly (full details are available in the Unit Information of this unit's Moodle site).

After connection you will see:



SQL Statements can be entered in the right hand panel, labelled "Worksheet".

2. Changing your password

After logging in for the first time you should change your password from the one supplied by your tutor.

DO NOT set your Oracle password the same as your standard authenticating password.

Oracle has several **important** limits on the password you set:

- The password is not case sensitive
- May be 1 - 30 characters in length
- Must begin with an alphabetic character
- Can contain only alphanumeric characters and the underscore, @, #, \$, and %

You should set your password in SQL DEVELOPER using the PASSWORD command (your password is hidden). Type the word password in your worksheet and click on the Execute Statement button (the green arrow) in the toolbar to run the command:



After changing your password for the first time please logout and then re log back in to check that your password change has been successful.

If at any stage you find that you are unable to login to Oracle due to password problems *please email your tutor to have your password reset.*

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder


1.2 Using SQL Developer GUI to manage data

There are two main approaches to interact with Oracle database using SQL Developer, through the Graphical User Interface (GUI) and SQL Worksheet. In this module, you will learn to use the SQL worksheet to create the database and to use the GUI to add, update and delete data from the database.

1. Opening an SQL file in the SQL Worksheet.

We will use SQL worksheet to create the database. Follow the these steps>

1. Download the myfirstDB.sql from tutorial 1 resources in Moodle.
2. Open the SQL Developer in your machine.
3. Open the connection to an Oracle server.
4. Drag the myfirstDB.sql to the SQL Worksheet area.



```

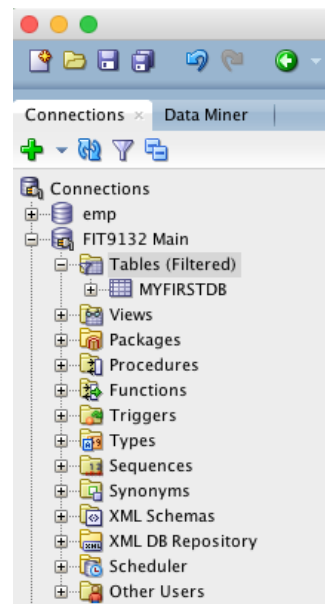
1  set echo on
2
3  create table myfirstdb
4      column1 number(5) not null,
5      column2 varchar(50) not null,
6      column3 varchar(50) not null,
7      constraint pk_myfirstdb primary key (column1));
8
9  -- Adding data to the table
10
11
12 insert into myfirstdb values (1,'a','aa');
13 insert into myfirstdb values (2,'b','bb');
14 insert into myfirstdb values (3,'c','bb');
15 insert into myfirstdb values (4,'d','dd');
16
17 COMMIT;
18
19 set echo off
20
21

```

5. Run the SQL script by pressing the "run script" icon. 

2. Viewing table structure and data inside a table.

1. To view the table using the graphical user interface, expand the Table option in the




Connection tab and find "myfirstdb" from the list. Double click on the "myfirstdb".

- You will see the listing of the database structure of the "myfirstdb" database.

COLUMN_NAME	DATA_TYPE	NULLABLE	DATA_DEFAULT	COLUMN_ID	COMMENTS
1 COLUMN1	NUMBER(8,0)	No	(null)	1 (null)	
2 COLUMN2	VARCHAR2(50 BYTE)	No	(null)	2 (null)	
3 COLUMN3	VARCHAR2(50 BYTE)	No	(null)	3 (null)	

- To view and change the data inside the myfirstdb, click on the "Data" tab on the right hand



- You will see the data listing inside the "myfirstdb".

	COLUMN1	COLUMN2	COLUMN3
1	1	a	aa
2	2	b	bb
3	3	c	bb
4	4	d	dd

3. Updating the data inside a table.

- Double click on the cell where the data needs to be changed. For example, let's change the content

	COLUMN1	COLUMN2	COLUMN3
1	1	a	aa
2	2	b	bb
3	3	c	bb
4	4	d	dd


of row-1 of column2 from "a" to "z".


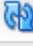




- You will see an asterisk for the first row. To accept the changes made on the row(s) with asterisk,

you will need to issue a COMMIT by pressing the "tick"  icon.







- When database accept the COMMIT instruction you have made, the asterisk should disappear from row 1.

4. Adding a new row to the table.

1. Click on the insert new row icon .
2. You will see a new row is being added to the table with all values listed as (null).

							Sort..	Filter:
	COLUMN1	COLUMN2	COLUMN3					
1	1	z	aa					
2	2	b	bb					
3	3	c	bb					
4	4	d	dd					
+5	(null)	(null)	(null)					

3. Replace the (null) value with 5, e and ee.
4. Click on the COMMIT icon to save the changes.


Columns	Data	Model	Constraints	Grants	Status
					
Sort..			Filter: <input type="text"/>		
	COLUMN1	COLUMN2	COLUMN3		
1	1	z	aa		
2	2	b	bb		
3	3	c	bb		
4	4	d	dd		
5	5	e	ee		

5. You will see the new being added to the table.

Assignment Project Exam Help

<https://powcoder.com>

5. Deleting a row from the table.

1. Click on the row to be deleted, for example, choose row 2.
2. Once the row is highlighted, choose the delete icon  from the menu.
3. You will see the row to be marked with negative sign at the front of the row number.

	COLUMN1	COLUMN2	COLUMN3					
1	1	z	aa					
-2	2	b	bb					
3	3	c	bb					
4	4	d	dd					
5	5	e	ee					

4. Click on the COMMIT icon to save the deletion.
5. The second row now should contain the data 3, c, bb.

1.3 Data Anomalies

A poor database design will lead to problems during the operation of the database. The following exercises will help you to identify anomalies(problems) associated with a poorly designed database.

First, you will need to create the database and to populate the data. To perform this task, you need to make sure you have completed the module Introduction to SQL Developer.

Creating a database

1. Download the file student_poor.sql and student_good.sql from Moodle.
2. Open the SQL Developer software.
3. Connect to the Oracle server.
4. Open the files you have downloaded from Moodle in SQL Developer.
5. Run the student_poor.sql.
6. Run the student_good.sql.

Manipulating data inside the database

There are two databases you have created after you completing steps 5 and 6.

1. The poorly designed database contains a single table called student_poor.
2. The well-designed database contains three tables called
 - student_good
 - unit_good
 - enrolment_good

Tasks

- Locate the listing of these tables in SQL developer to check whether you have completed step 5 and 6 of Creating database correctly.
- Performs the database operations instructed below and observe the impact of the action on the data. Several questions have been provided to help your observation.
- For each of the action, perform the action in both the poorly designed and well designed databases. Compare the observations when the action took place in those databases.

DATABASE OPERATIONS

UPDATE

1. Change the name of FIT9131 for student number 1111 into Foundation of Java in the student_poor table.
2. Observe the data in the student_poor table.
 - Can you make the change?
 - How many different names does FIT9131 have? What will be a potential problem with this situation?
 - If you want to ensure the name of FIT9131 to be consistent, how many rows did you suppose

- to change in the database? Is it possible for you to easily check that all relevant rows have been updated?
- Now, perform the same update in the well design database.
 - What table do you need to change to reflect the required change?
 - How many row(s) do you need to change?
 - How does this design ensure consistency of the unit name?
 - What problem(s) does the well design database overcome?

DELETE

- Wendy Wang decided to withdraw from FIT9133. You are required to delete this enrolment from student_poor table.
- Observe the data in the student_poor table.
 - Can you find out from the database the name of a unit with the code FIT9133 after the delete operation is completed?
 - Can you find out from the database the student number of Wendy Wang after the delete operation is completed?
- Now, perform the same deletion in the well design database.
 - In what table the data need to be deleted?
 - Can you find the student number of Wendy after the deletion?
 - Can you find the unit name for FIT9133 after the deletion?

Assignment Project Exam Help

INSERT

- A new unit FIT5000 Data Analytics is introduced in the course. Add the new unit code and unit name to the student_poor table.
- Observe the data in the student_poor table.
 - Can you perform the insert operation? What error message did you get?
 - What additional information do you need to include to add a new unit information to the student_poor table? Is it a good idea? Why/Why not?
- Now, perform the same insertion in the well design database.
 - In what table the data need to be added?
 - Can you add the new unit information in the database without having any student enrolled into the unit?

Summary

Designing what data and how it is stored in an organisation is important. A poorly design data store may lead to problems such as those you have seen in the exercises. The poor_student table used in this module contains three different things/objects (physical object - student and abstract concepts - unit, enrolment). These three objects/concepts are separated into three different tables in the well design database. This is how data should be typically stored in the database. In this unit, you will learn how to ensure that you will build database without the associated problems shown in this exercise. This will be done by looking at Relational model, its theory and design methodology.

2

Week 2: The Relational Data Model FIT9132

In this week's work, we will look at the fundamental concepts on which the Relational Database model is built.

Part A: The Relational Model

1. Discuss the following terms

- Relation
- Attribute
- Domain
- Tuple
- Degree and Cardinality of a Relation

2. Consider the CUSTOMER and ORDER relations below:

CUSTOMER (CUSTOMER-ID, NAME, ADDRESS)
ORDER (ORDER-ID, DATE, CUSTOMER-ID)

Assume a single customer may have any number of orders.

1. Identify the primary key and foreign key attributes in these relations.
2. Can you think of a reason why we would not just store all the customer and order information in one relation so that we would not have to perform the join operation?

3. Choosing Primary key

1. In any relation, tuples must be unique. However, in many cases, the set of all the attributes in a relation is not considered a candidate key. Why not? On the other hand, suppose we do have a relation where the set of all attributes is a candidate key. In this case, show that this set must, therefore, be the only candidate key and hence the primary key.
2. Consider a relation that depicts a tutorial room booking in a university. Each faculty assigns a person to handle the booking for all tutorial classes for that faculty. The person's email address is given to the university's booking system as a contact person. BOOKING(b_date, b_starttime, b_endtime, unit_code, contact_person, room_no, tutor_id)
 - Identify candidate key(s) and primary key for the relation if the following business rules are applicable:
 - More than one tutorial classes of the same unit may run at the same time (parallel sessions are possible).
 - A tutor may teach several classes of the same unit.
 - All tutorial classes are 2 hours long.
 - Identify candidate key(s) and primary key for the relation if the following business rules are applicable:
 - Tutorial classes can be either 1 hour or 2 hours long.
 - A tutor can only teach one tutorial class in a given unit.
 - There are no parallel sessions of tutorial classes.

4. Relational Algebra (Adapted from Exercise 3.6 of Connolly, Begg and Strachan)

Suppose we have the following 4 relations:

HOTEL(HOTEL-NO, NAME, ADDRESS)
 ROOM(ROOM-NO, HOTEL-NO, TYPE, PRICE)
 BOOKING(HOTEL-NO, GUEST-NO, DATE-FROM, DATE-TO, ROOM-NO)
 GUEST(GUEST-NO, NAME, ADDRESS)

Generate the relational algebra for the following queries:

1. List the names and addresses of all hotels
2. List all single rooms with a price below \$50
3. List the names and addresses of all guests
4. List the price and type of all rooms at the Grosvenor Hotel
5. List all names and addresses of guests currently staying at the Grosvenor Hotel (assume that if the guest has a tuple in the BOOKING relation, then they are currently staying in the hotel)

The following question is an optional exercise.

Assignment Project Exam Help

5. There are 7 relational algebra operators, namely:

SELECTION, PROJECTION, JOIN, UNION, INTERSECTION, DIFFERENCE and CARTESIAN PRODUCT. In fact, 5 of these operators may be considered primitive operators in the sense that the others may be expressed in terms of the primitive operators. The primitive operators are:

SELECTION, PROJECTION, JOIN, UNION, INTERSECTION, DIFFERENCE and CARTESIAN PRODUCT

Using the sample tables below, show how the JOIN operation can be expressed in terms of the fundamental operators by showing the process to do a natural join of customer and order.

- **CUSTOMER table:**

Cust_ID	Name
1	Green
2	Blue

- **ORDER table:**

Ord_ID	Date	Cust_ID
1	23-Feb-2009	1
2	26-Feb-2009	1
3	26-Feb-2009	2

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder


3

Week 3: Conceptual Modelling

Conceptual Design - Theory Questions

Reference: Coronel, C and Morris, Database Systems: Design, Implementation & Management, Chapter 4, Selected Review Questions and Problems.

1. What two conditions must be met before an entity can be classified as a weak entity? Give an example of a weak entity.
2. What is a strong (or identifying) relationship, and how is it depicted in a Crow's Foot ERD?
3. Given the business rule "an employee may have many degrees," discuss its effect on attributes, entities, and relationships. (Hint: Remember what a multivalued attribute is and how it might be implemented.)
4. What is a composite entity, and when is it used?
5. Suppose you are working within the framework of the conceptual model in Figure Q4.5. Figure Q4.5 The Conceptual Model for Question 5



- o Write the business rules that are reflected in it.
 - o Identify all the cardinalities.
6. What is a recursive relationship? Give an example.
 7. How would you (graphically) identify each of the following ER model components?
 - o an entity
 - o the cardinality (0,N)
 - o a weak relationship, and
 - o a strong relationship
 8. Discuss the difference between a composite key and a composite attribute. How would each be indicated in an ERD?
 9. What two courses of action are available to a designer when he or she encounters a multivalued

attribute?

10. What is a derived attribute? Give an example.
 11. How is a relationship between entities indicated in an ERD? Give an example, using the Crow's Foot notation.
 12. Discuss two ways in which the 1:M relationship between COURSE and CLASS can be implemented. (Hint: Think about relationship strength.)
 13. How is a composite entity represented in an ERD, and what is its function? Illustrate using the Crow's Foot model.
 14. Briefly, but precisely, explain the difference between single-valued attributes and simple attributes. Give an example of each.
 15. What are multivalued attributes, and how can they be handled within the database design?
-

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

3.1 Using Tools to draw ERD

Tools to draw ER diagrams

There are several tools available to draw ER diagrams. Some of them are available to be used within a web browser. Some examples of these are:

- **Lucidchart** (<https://www.lucidchart.com/>) - this product is a browser-based diagramming tool; it is able to draw a wide range of different diagrams, including ER Diagrams. As a University student, you are entitled to a free Lucidchart account. Your account has been created and an invitation email sent to your Monash student account (check Spam if the message is not in your inBox). Click on the link in your Lucidchart welcome email and you will be able to make use of the software. As a first step, you should look at the provided tutorials, in particular, "[Entity Relationship Diagrams](https://lucidchart.zendesk.com/entries/21606135-Entity-relationship-diagrams)" (<https://lucidchart.zendesk.com/entries/21606135-Entity-relationship-diagrams>)" (begin with "Manual ERD Creation", note that the model we are asking you to build should not include the "Type" column), or

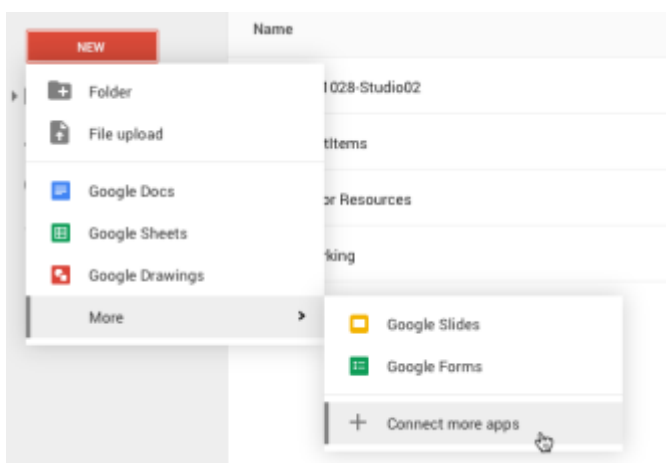
- any other drawing package you wish. One excellent alternative is Gliffy (<https://www.gliffy.com/>)

At this stage of our study, we do not wish to use a CASE tool - it is important that we first establish a clear understanding of Entity Relationship modelling.

Setting up Lucidchart

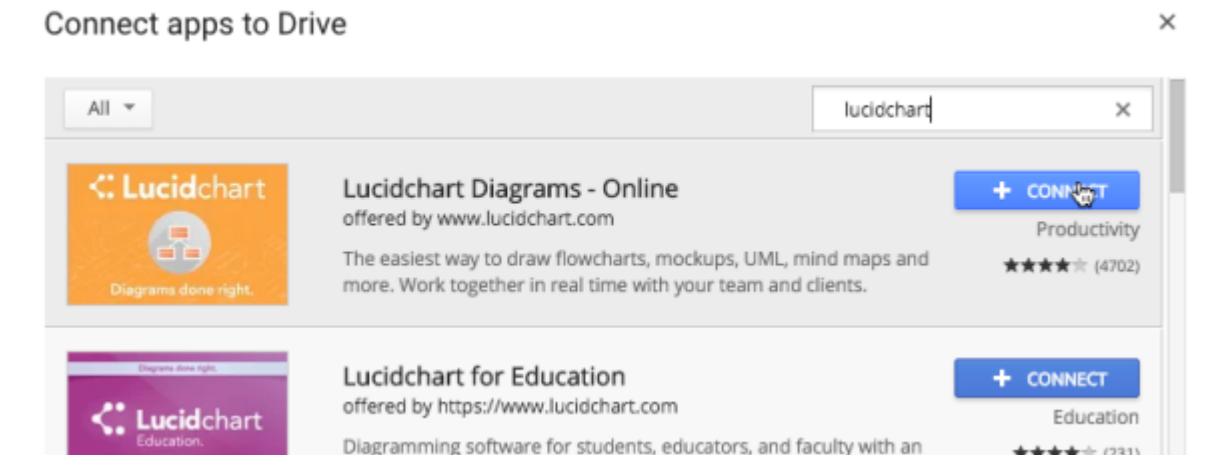
After you have setup your account, Lucidchart can be logged into directly at the Lucidchart site: <https://www.lucidchart.com/> or better, accessed from your student Google Drive (from my.monash select the "Drive" link).

Connect to Google Drive, click "New" and then more, and "Connect more apps"



This will open the app picker which lists a wide range of applications that can be added to your Google drive.

You can browse through and find Lucidchart or alternatively type "Lucidchart" in the top right search box:



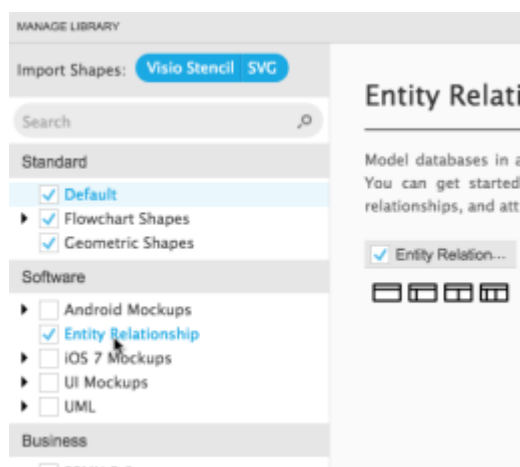
Select the "+ CONNECT" button for Lucidchart Diagrams - Online (do not select Lucidchart for Education)

After this process has been completed Lucidchart charts created from within your Google Drive will be stored in your Google Drive and not on the Lucidchart server.


To create a new Lucidchart - simply select New in Google Drive, and then select "Lucidchart" (you may need to expand "More" to see the Lucidchart option).

If this is the first document you have created as a Lucid Chart from your Google Drive you will be asked to approve Lucidchart's access to your Google Drive - please ensure you select "Accept". This will then transfer you into the Lucidchart workspace with a new document open. Select the "Entity Relationship (ERD)" template, then Blank ERD and then "Start Drawing".

If the ERD shapes are not listed in the left panel, add the ERD shapes by selecting "More Shapes" (bottom left) and then checking "Entity Relationship" and then "Save":



The symbol we will use to represent an entity is the second symbol from the left:

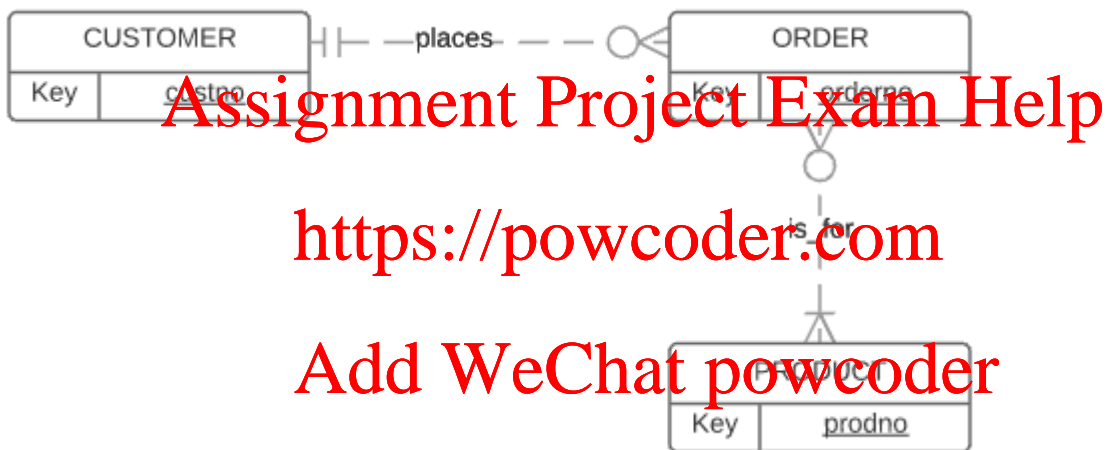


Using Lucidchart

Given a scenario represented by the following entities, where customers place orders for products:

- CUSTOMER - customer number, name, address, phone number
- ORDER - order number, order date and for each product ordered the quantity ordered and the total line price
- PRODUCT - product number, product description and product unit price

An initial ERD using Lucidchart for this scenario would be:



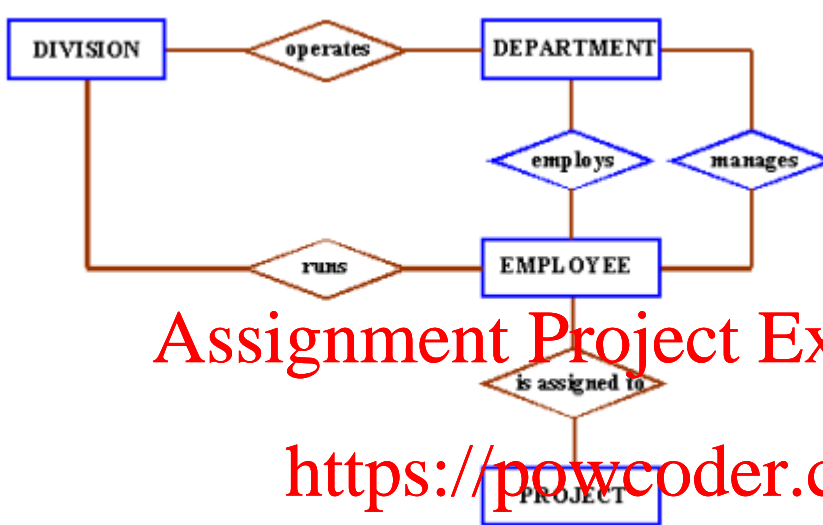
This ERD only shows the keys of each of the entities, sometimes an ERD is drawn such that it will show all the non-key attributes for the model as well. In such a complete ERD you must not show or label foreign keys, the use of foreign keys indicates that you are looking at a logical model where a choice has been made to use a relational database, rather than a conceptual model.

Prepare the ERD shown above using your choice of drawing tool and then create a copy of this initial model and add all of the non-key attributes listed in the scenario above to your copy. In Lucidchart, you can easily create a copy by selecting the bottom tab of the diagram, click the down arrow and then select "Duplicate".

3.2 Building Conceptual Models

Conceptual Modelling - Practical Work

For this weeks lab exercises we are going to prepare conceptual models (Entity Relationship Diagrams - ERDs) for a number of scenarios. At this point we are not interested in the database implementation of these models, our aim will be to model without any consideration of the database system in which the model may ultimately be implemented.



1. Use the above diagram and the business rules below, create a Entity Relationship Diagram using crow's-foot notations.

- Include:
 - all appropriate connectivities,
 - all cardinalities and
 - at least the minimum number of attributes required to implement the model
- Business Rules:
 - A department employs many employees, but each employee is employed by one department.
 - Some employees, known as "rovers," are not assigned to any department.
 - A division operates many departments, but each department is operated by one division
 - An employee may be assigned to many projects, and a project may have many employees assigned to it.
 - A project must have at least one employee assigned to it.
 - One of the employees manages each department, and each department is managed by one employee
 - One of the employees runs each division, and each division is run by one employee.

2. Prepare an Entity Relationship Diagram (ERD) showing all attributes and the identifier of each entity for the following description of a Property Rental System:

- Properties are rented by tenants. Each tenant is assigned a unique number by the Agency. Data

held about tenants include family name, given name, property rented, contact address - street, city, state, postcode & telephone number. A tenant may rent more than one property and many tenants may rent parts of the same property (eg. a large shopping complex).

- Properties are owned by owners. Each property is assigned a unique building number. The agency only recognises a single owner for any of the properties it handles. The owner, address, and value are recorded for each property. In addition the lease period and bond are recorded for each property or sub property rented. An owner may own several properties.
 - Properties are subject to damage and the agency records all instance of damage to its properties - property, date, type of damage and repair cost are recorded. Repair costs are charged directly to tenants
 - Normal property maintenance is also noted - property, date, type of maintenance and cost are recorded. Maintenance costs are charged to the property owner.
 - Tenants pay accounts to the Agency - these consist of weekly rental payments, bond payments (for new properties) and damage bills. The date of payment, tenant, property, type of account (Rental, Bond, Damage) and amount are recorded. Each payment is assigned a payment number.
-

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

4 Week 4: Logical Modelling

Logical Modeling

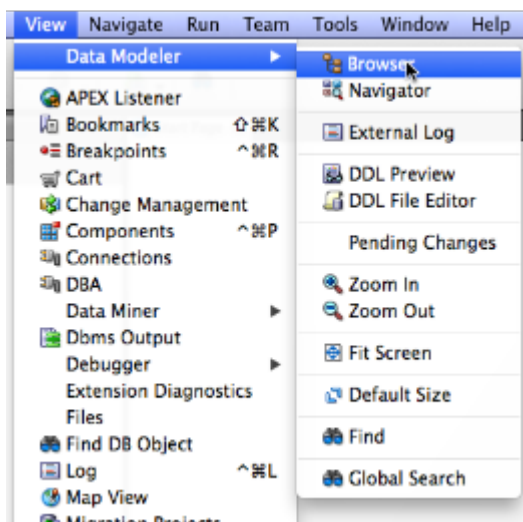
After preparing a conceptual model the next stage is to select the type of database we will use (for us relational) and convert our conceptual model into an appropriate logical model. For logical modelling we will make use of Oracle SQL Developer Data Modeler.

Task A: Using SQL Developer Data Modeler


This software is a commercial level tool with an extensive set of features including support for Subversion versioning and source control thus permitting teams of developers to work collectively on a design. Given the extensive range of features we are only going to be looking at a subset of these for our study. SQL Developer Data Modeler begins with a (relational) LOGICAL model and then creates what it calls a RELATIONAL model from which the schema file can be generated. The relational model is essentially a graphical representation of the physical model.

Accessing Data Modeler

Data Modeler is installed as part of the standard SQL Developer installation. To work with Data Modeler it is helpful to have the Data Modeler browser open in the left panel of SQL Developer. To open this panel select View - Data Modeler - Browser.



The browser opens with a new (unnamed) model:




To begin creating a logical model, right click Logical Model and select Show. This will open a blank model in the main working panel of SQL Developer and add a range of new icons to the main toolbar. Hover over each of these new icons to become familiar with what they represent. You should regularly save your design - right click on the design (here Untitled_1 as the design has not been saved as yet) and select a save location, after the first save the design will be named.

Assignment Project Exam Help

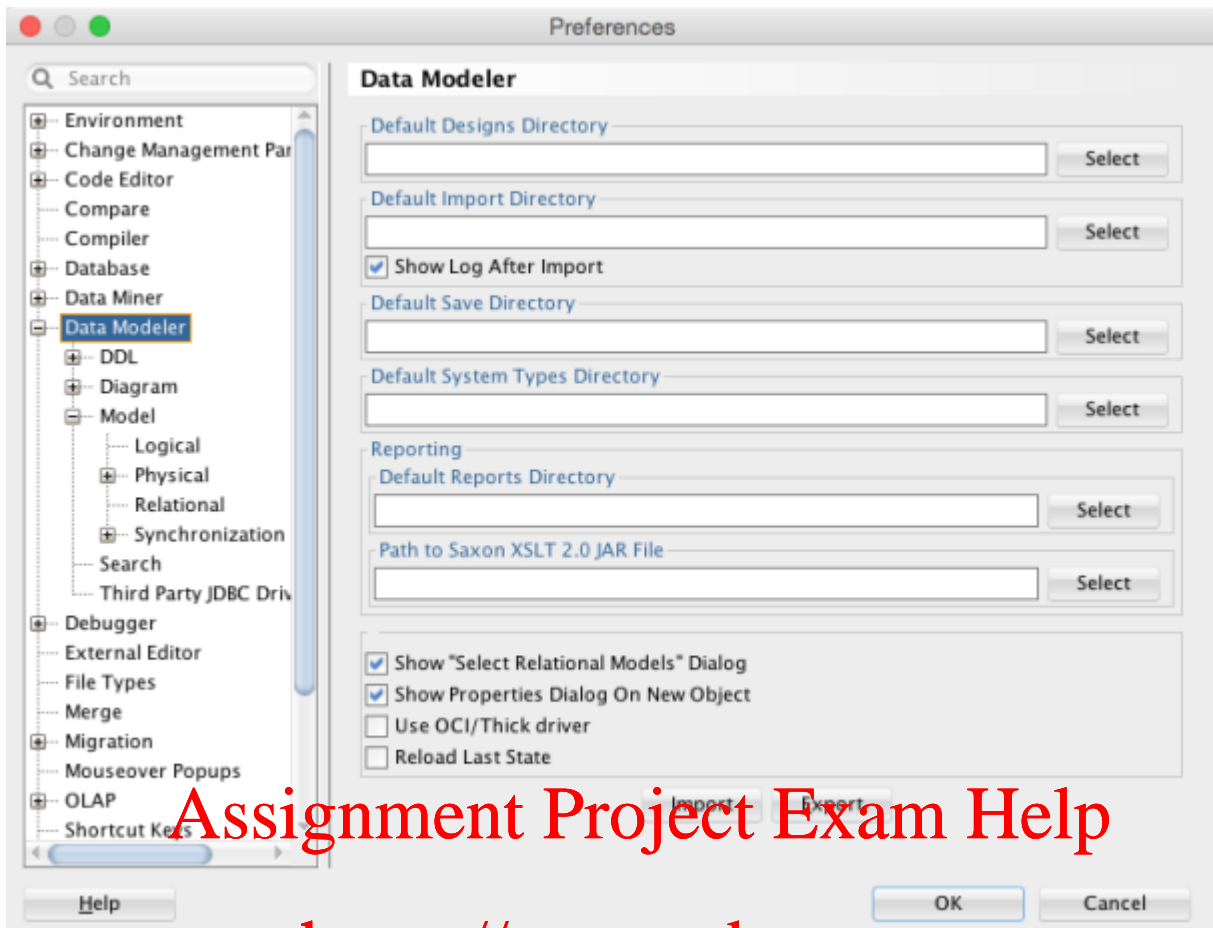
When using Data Modeler it is very important that you save and close your model before exiting SQL Developer or shutting your laptop. Failure to do so may result in loss of parts of your model.

<https://powcoder.com>



Add WeChat powcoder


A number of features should be configured using the SQL Developer preferences (Windows: Tools-Preferences, OSX: Oracle SQL Developer-Preferences), select the Data Modeler:



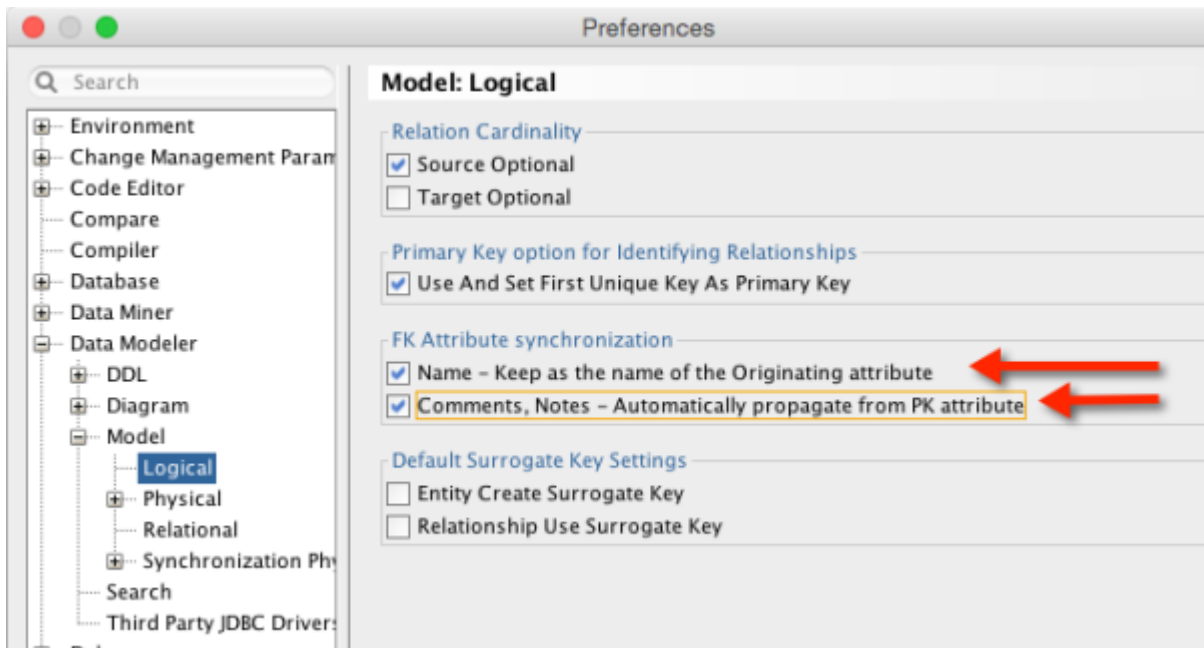
<https://powcoder.com>

We wish to modify two features:

- (a) select DDL and select the option to "Generate Short Form of NOT NULL Constraint" - this will cause not null constraints to be not named



- (b) select Model, Logical and check both items in "FK Attribute synchronization"



There are a large number of other settings available to configure, you might like to investigate these if you are using your own copy of SQL Developer, for the labs we will leave the remaining settings at the default values.

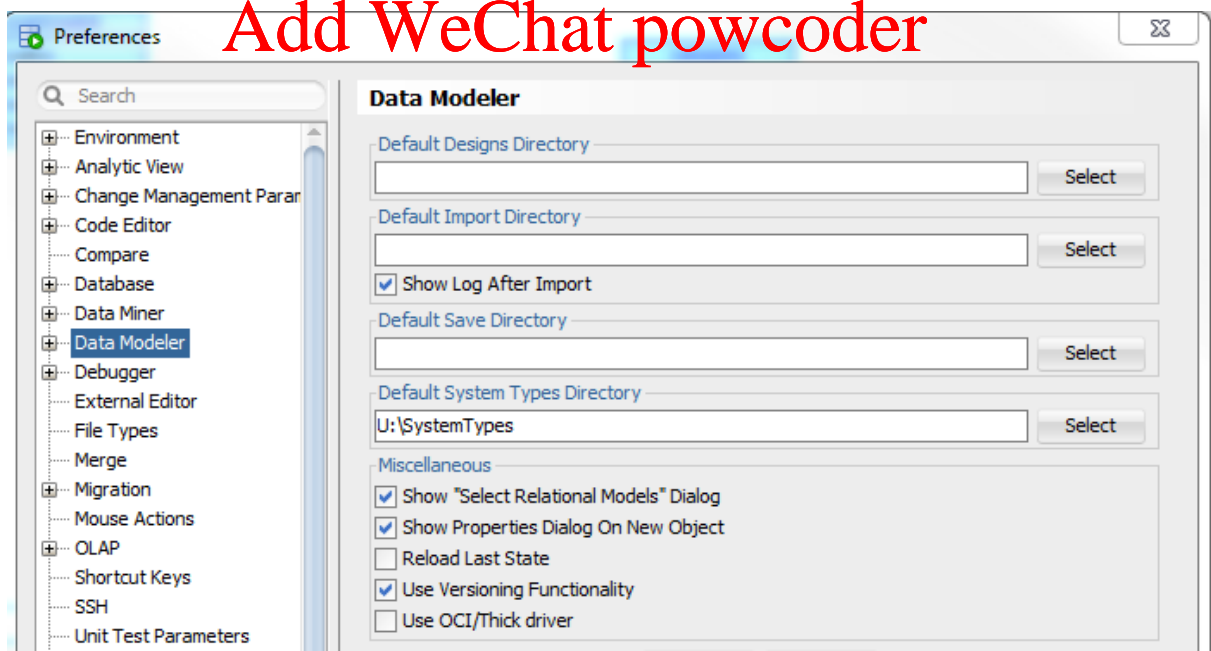
Assignment Project Exam Help

If you are working on a **Monash on-campus PC** please set up a folder to save the Modeler System Files:

<https://powcoder.com>

- In file manager, connect to your mapped drive (U:) and create a folder called SystemTypes
- In the Data Modeler preferences point the "Default System Types Directory" to this newly created folder

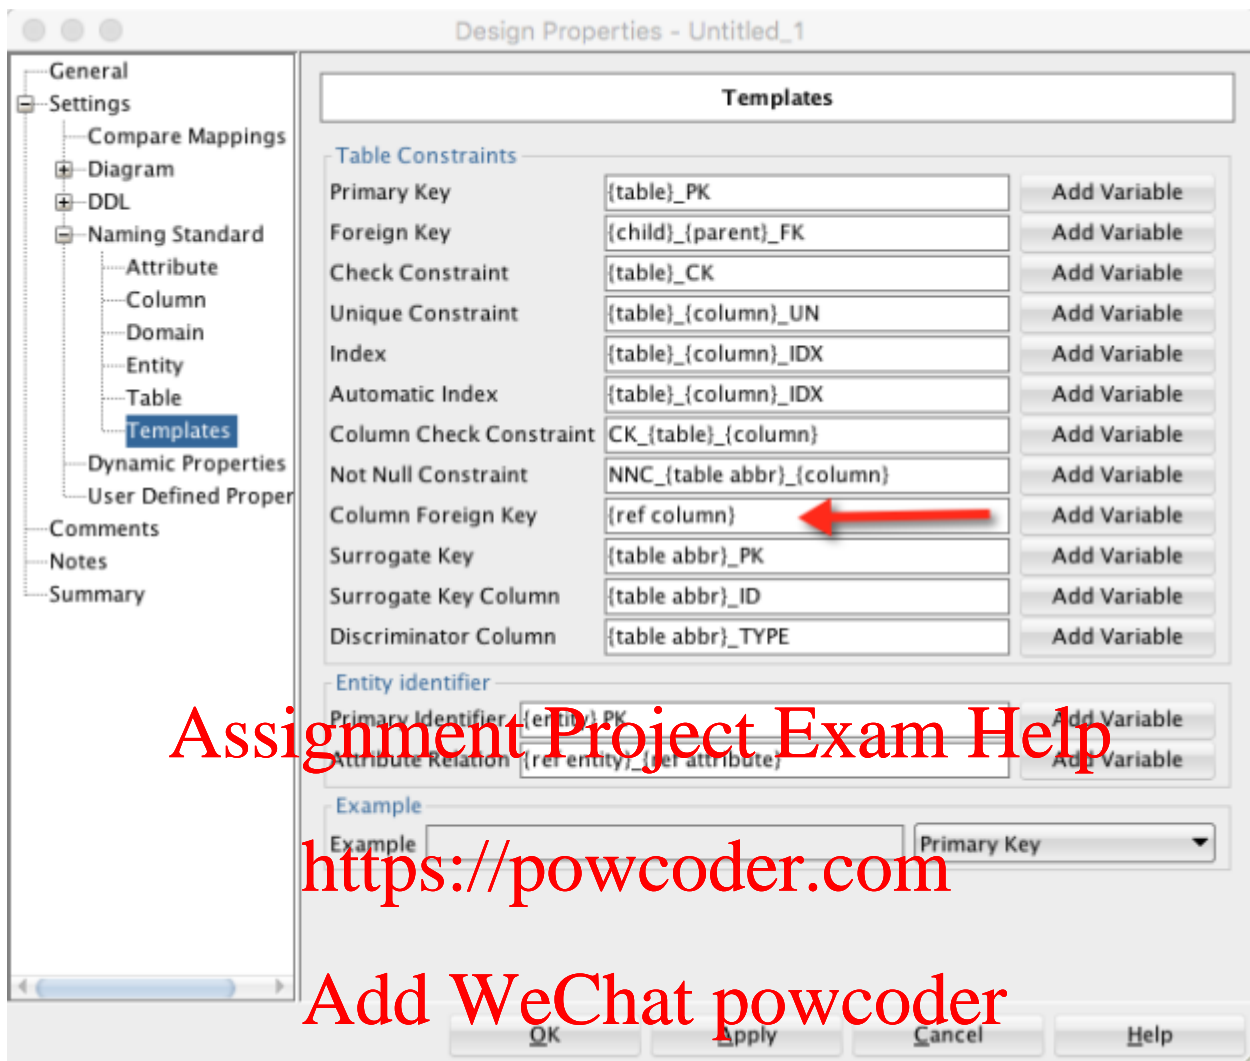
Add WeChat powcoder



Develop a model - Stage 1 The Logical Model:

Before starting a new project, right click the project in the Data Modeler browser and select properties. Within the Design Properties sheet select Settings - Naming Standard - Templates and modify the

"Column Foreign Key" setting from {ref table}_{ref column} to {ref column}:



For your first model we will implement a Customer-Orders system, represented by the following entities, where customers place orders for products:

CUSTOMER - customer number, name, address, phone number

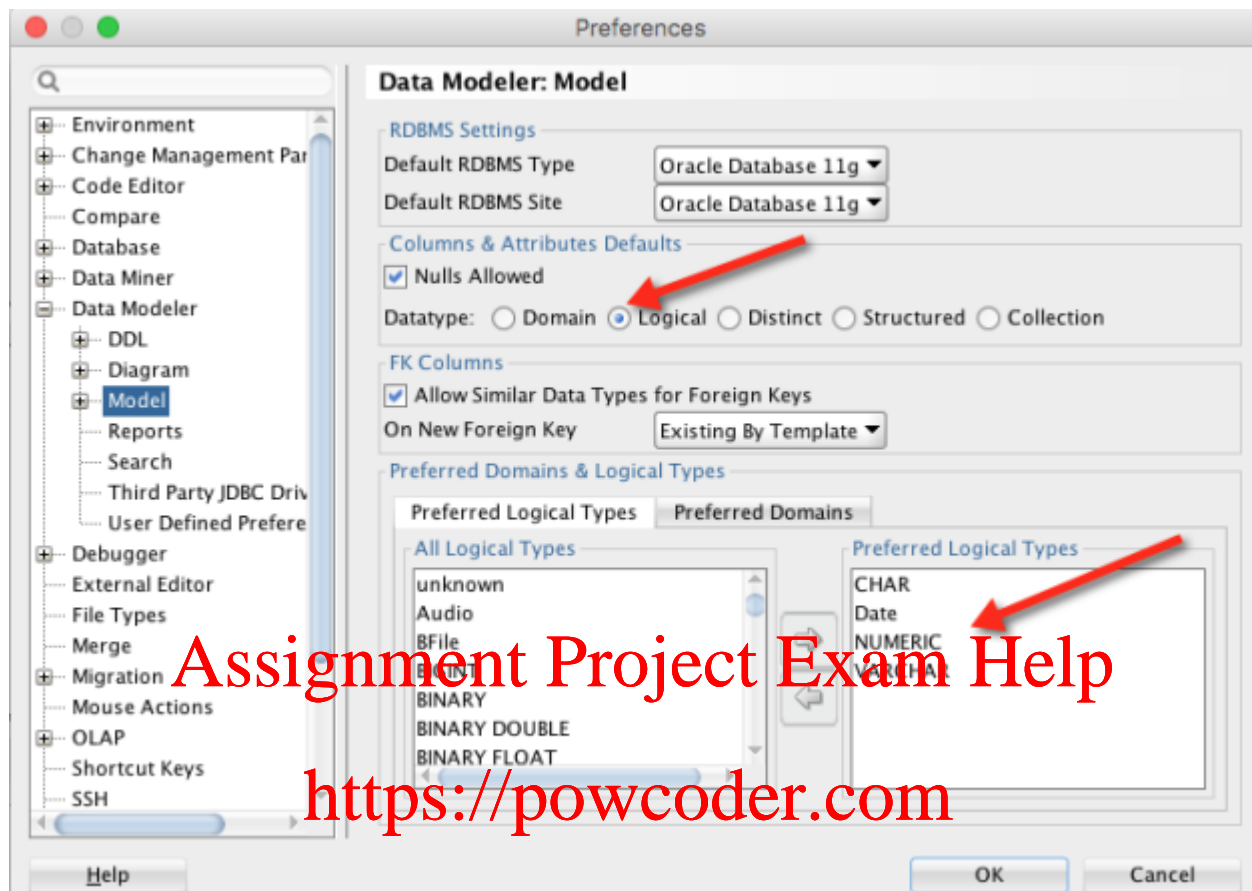
ORDER - order number, order date, customer number and for each product ordered the quantity ordered and the total line price

PRODUCT - product number, product description and product unit price

Attributes on the logical model have a number of possible data types, the main ones for us being "Domain" and "Logical":

- Domain types are domains which you create via the menu items - Tools - Data Modeler - Domains Administration. No domain types are supplied, you must create any you need
- Logical types are not actual data types - they are names which are mapped to native types at a later stage. These logical types are pre populated with several Oracle types

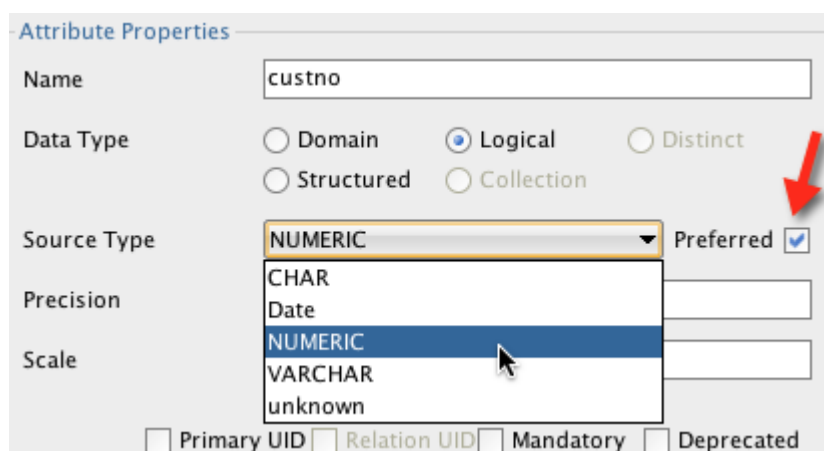
For our work we will not make use of domain types, instead you should always use logical types. You can speed up entry of attributes by restricting Data Modeler to logical and types and a preferred range of types. In your SQL Developer preferences set:



Add WeChat powcoder

The 'Preferred Logical Types' are populated by selecting the item in the 'All Logical Types' and clicking on the right pointing arrow between All and Preferred.

Then when adding an attribute:



check the 'Preferred' tick box and you will only see those types you have selected as preferred (clearly you may modify the set preferred data types to suit your needs).

On your logical model add an entity named CUSTOMER (name the entity under the General features in

the Entity Properties dialog box) and then select the attributes feature. Add the following attributes:

- custno - Logical type: Numeric Precision 7, Scale 0, Primary UID
- custname - Logical type: VARCHAR(50), Mandatory
- custaddress - Logical type: VARCHAR(50), Mandatory
- custphone - Logical type: CHAR(10)


For each attribute add a meaningful description of the attribute under the attribute option - "Comments in the RDBMS". Then add an ORDERS entity with attributes:

- orderno - Logical type: Numeric Precision 7, Scale 0, Primary UID
- orderdate - Logical type: DATE, Mandatory

Note here that the entity is being named as ORDERS since ORDER is a reserved SQL word - remember that as a general rule entity names must not be pluralised.

For all the logical models we create you should set (right click on the logical model) the

- Notation to "Information Engineering Notation",



- View Details to "Attributes", and
- Show to "Labels" and "Legend"

When modelling students are required to include the "Legend" on all models (the panel may be moved to fit your models layout). **Models submitted without a legend will not be graded.**

Now add a 1:N Relationship between CUSTOMER and ORDERS - click in CUSTOMER (the parent) and then in ORDERS. In the General Relations Properties name the relation "places".

Enter a name for the source and target, and set up the participation (Customer - optional, Order - mandatory ie. not optional):

Relation Properties - places

General

Name:

Use surrogate keys: ☐

Source Cardinality

Source:

Source key:

Name on Source:

Source Entity Synonym:

Source to Target Cardinality:

Source Optional: ☒

Transferable: ☐

Target Cardinality

Target:

Target key:

Name on Target:

Target Entity Synonym:

Target to Source Cardinality:

Target Optional: ☐

Transferable: ☒

Dominant Role:

Identifying: ☒

Delete Rule:

Assignment Project Exam Help

<https://powcoder.com>


Add WeChat powcoder

Proceed and add the PRODUCT entity.

Modeler can draw M:N relations(hips) on its logical models, however you cannot add attributes to such relationships - ***please do not use the M:N relation(ship) in our models.***

Add a new entity ORDERLINE and then connect this with ORDERS and PRODUCT via 1:N Identifying relations(hips).

Your final logical model will have the form:

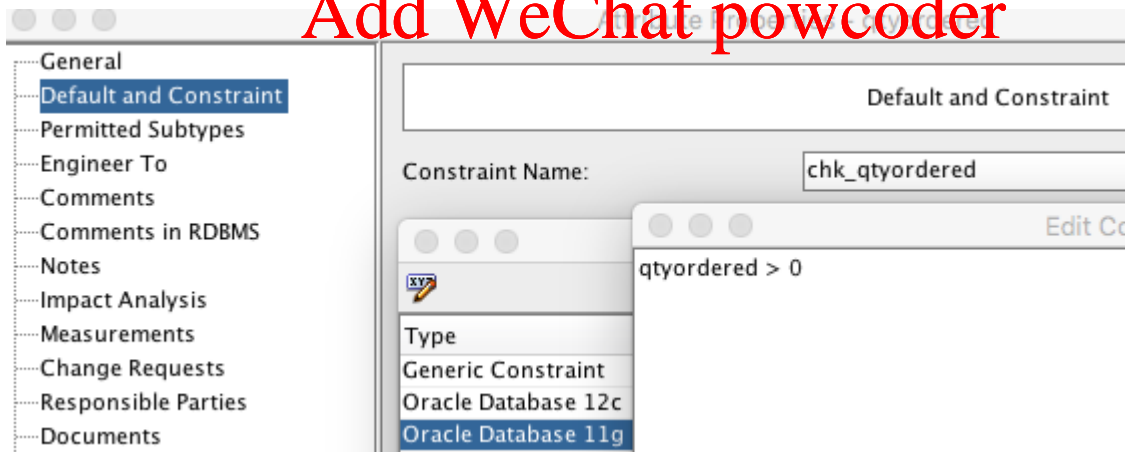


Assignment Project Exam Help

Constraints can be added to the logical model, as an example lets add a constraint to say that qtyordered in ORDERLINE must be greater than zero. Select the ORDERLINE entity, then select attributes and double click on qtyordered. In the left hand list of the pop-up window select "Default and Constraint". Give the constraint a name for example "chk_qtyordered" (be careful to select an informative name eg. chk_columnname) and then enter the constraint into the constraint editor:

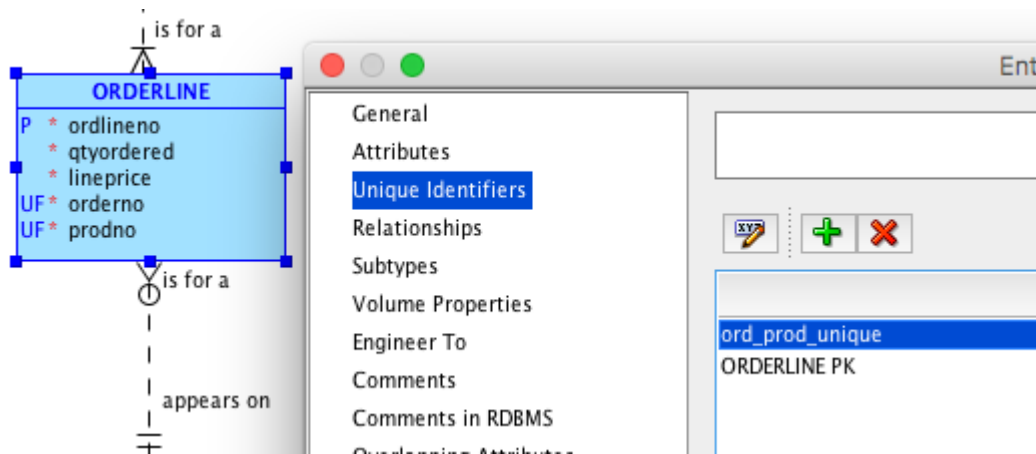
<https://powcoder.com>

Add WeChat powcoder



Note that the constraint added is the inner part of the standard SQL CHECK clause.

In developing a logical model, where *appropriate and documented*, you are free to introduce surrogate primary keys. If a surrogate PK is introduced into your design you must ensure that the original key's uniqueness is still maintained by enforcing a unique identifier which includes the attributes which compose the original key. As an example, say you decided to identify order lines by a unique order line number (*this would not be a good/common choice*, it is being used here for demonstration purposes only). On the ORDERLINE entity, you need to select Unique Identifiers and add a new identifier with the attributes in your previously identified natural key (here orderno and prodno):



Develop a model - Stage 2 The Relational (Physical) Model:

This completed logical model is now "Engineered" to a "Relational Model".

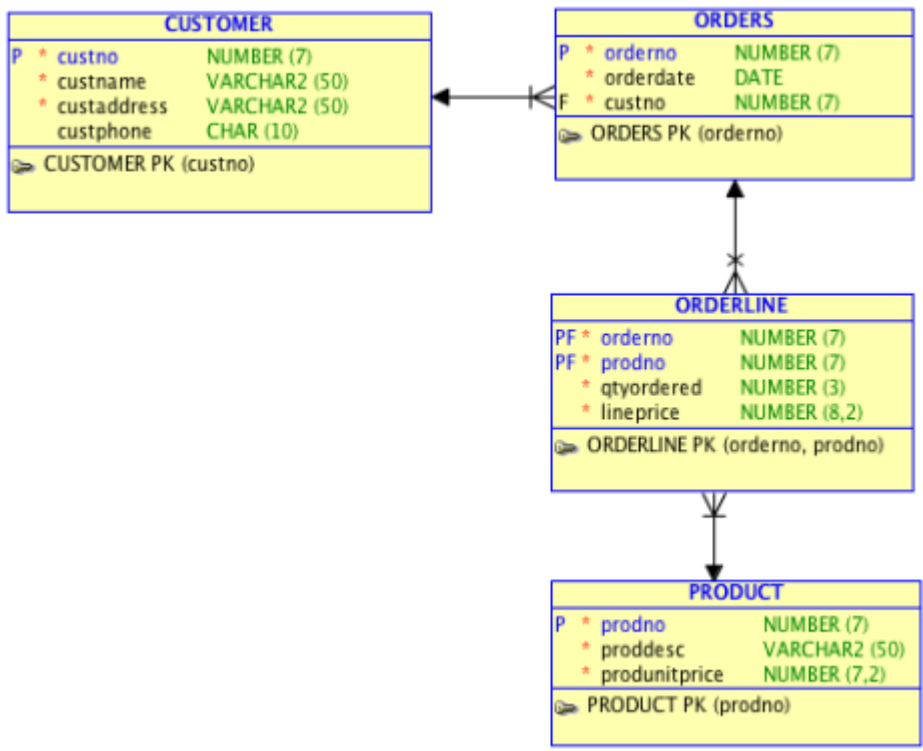


Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

In the pop-up window which appears, select "Engineer" the bottom left button. The logical model will then be engineered and a Relational Model is opened:



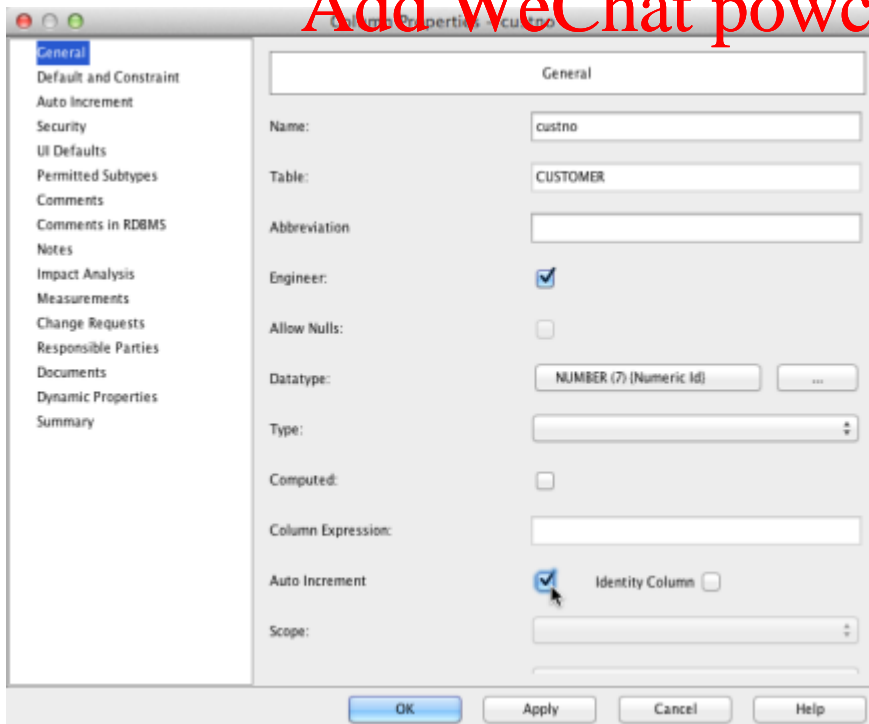
Assignment Project Exam Help

Note that in the default preferences setting, the Data Modeler relational model shows relationship lines pointing from the FK to the PK in the parent entity.

<https://powcoder.com>

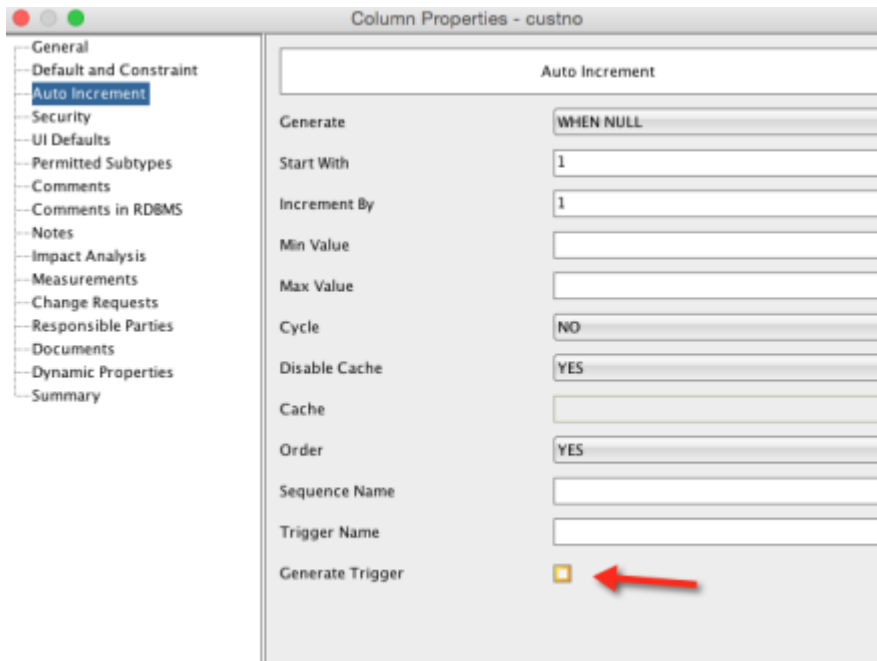
As an example of what can be configured in the Relational model, select the CUSTOMER table, then the columns, then double click on the custno column, in the dialog which opens select "Auto Increment"

Add WeChat powcoder



This will result in a sequence being created when the DDL is generated. The default configuration also generates a trigger to support the auto increment which we do not need. To prevent this trigger being

generated, after you have selected "Auto Increment" on General (above), select the Auto Increment option on the left and uncheck "Generate Trigger".

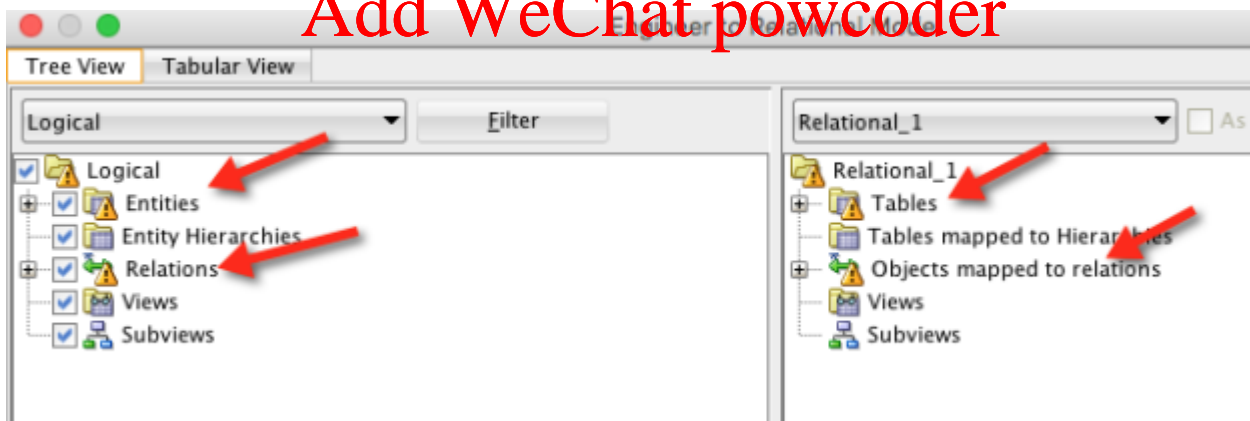


Assignment Project Exam Help

If you are re-engineering a model (ie. trying to generate a previous relational model, after changes to your logical model) it is very important that you note that SQL Developer does not automatically sync deletions from your logical model - such changes must be individually selected to be synced to your relational model. When re-engineering a previous model carefully check the "Engineer to Relational Model" for any triangles with an exclamation mark symbol:

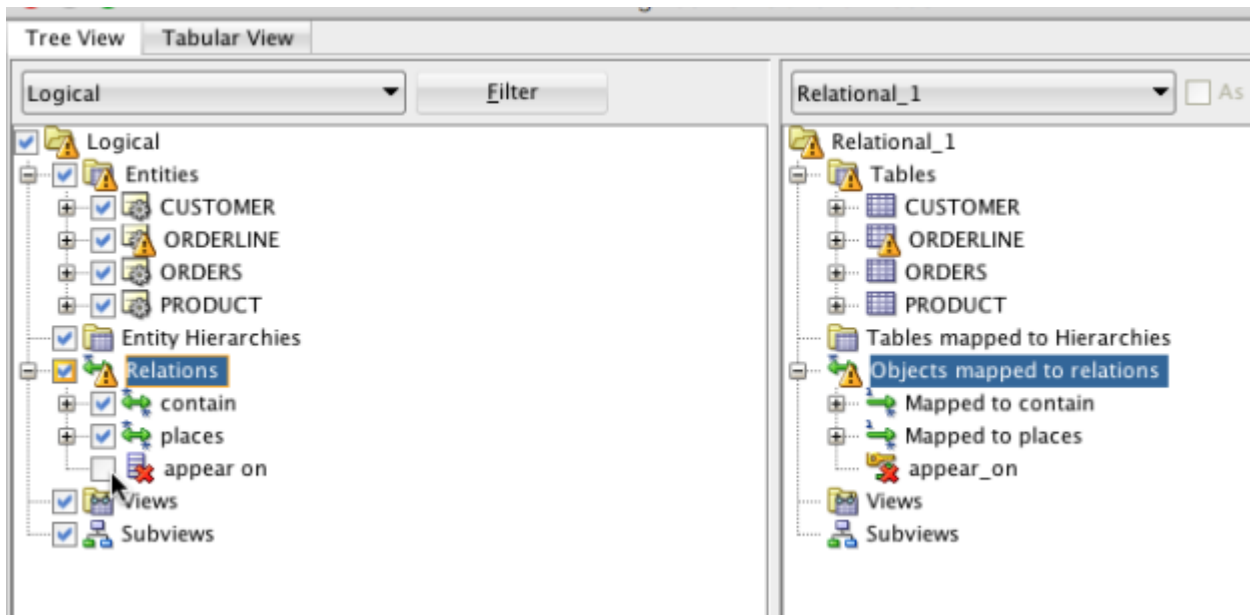
<https://powcoder.com>

Add WeChat powcoder



Such symbols represent issues you **must address** before generation.

For example here under relations (where we are removing the "appear on" relation as a demonstration of what occurs):




the removal of "appear on" has not been selected to be engineered to the relational model. You need to check the box if you wish it to be engineered (which we normally would).

If your relational model gets very confused you can select the relational model tab, do a ctrl+A or Apple+A and delete all the objects. The model can then be regenerated. Under **no circumstances should you delete the relational model itself** (in the left browser navigator), a bug in several versions of the software can result in such an action causing your relational model to "disappear".

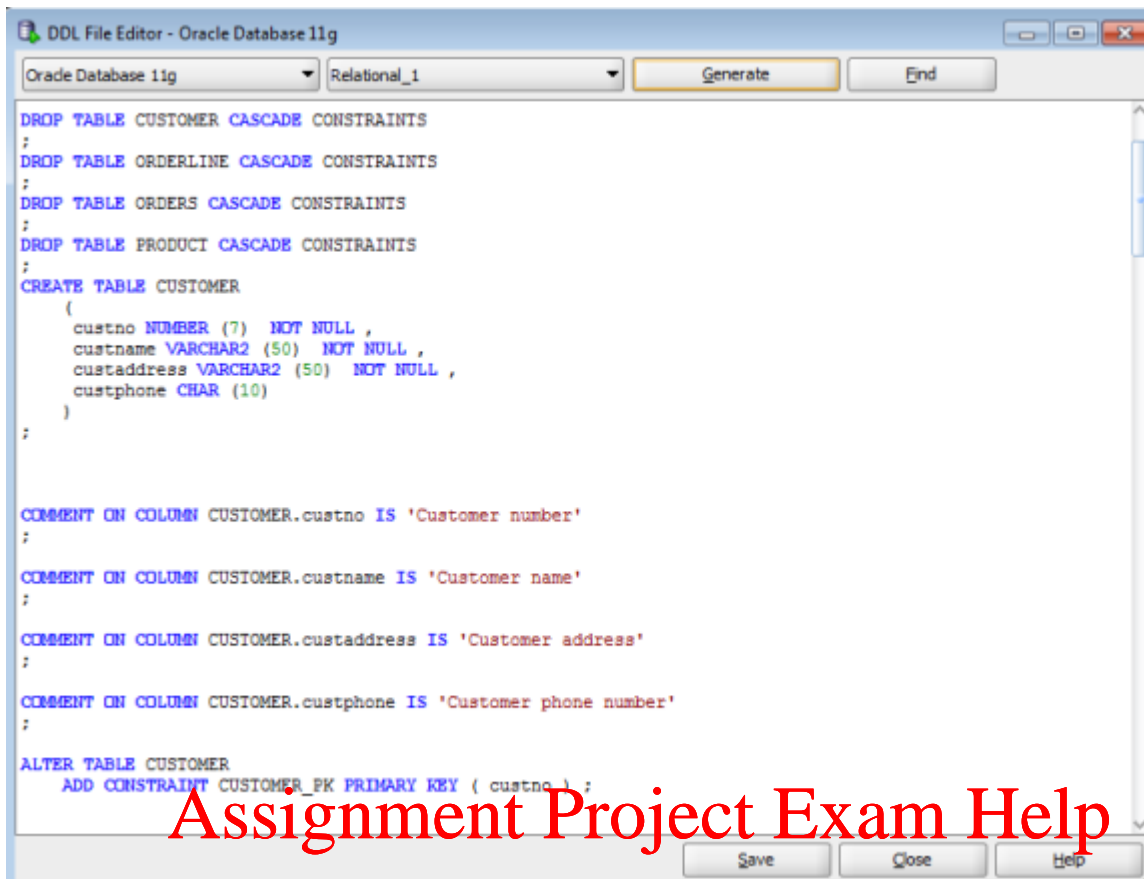
<https://powcoder.com>

When you have configured the relational model as you wish, select Generate DDL from the top toolbar:



Add WeChat powcoder

Select "Generate" in the pop-up window, specify the DDL Generation Options you wish (Drop tables should be included) and then select OK to generate the DDL.



Assignment Project Exam Help

<https://powcoder.com>

The generated file can be saved as an Oracle schema script (ensure you use an extension of .sql). Test your generated file against Oracle and confirm it operates correctly.

Add WeChat powcoder

It is also possible to configure Data Modeler to directly synchronise the design into the Data Dictionary of a database connection (we will not use this approach).

4.1

Logical Modelling - Task B - Rental Model

Task B: Using SQL Developer Data Modeler

Using your model from last week for the "Property Rental System", map your Conceptual model (ERD) into a logical model in Oracle SQL Developer Data Modeler.

Engineer your Logical Model to a Relational Model and then create the tables etc in Oracle from the generated DDL. In doing so make use of at least one check clause and one sequence.

- Properties are rented by tenants. Each tenant is assigned a unique number by the Agency. Data held about tenants include family name, given name, property rented, contact address - street, city, state, postcode & telephone number. A tenant may rent more than one property and many tenants may rent parts of the same property (eg. a large shopping complex).
- Properties are owned by owners. Each property is assigned a unique building number. The agency only recognises a single owner for any of the properties it handles. The owner, address, and value are recorded for each property. In addition the lease period and bond are recorded for each property or sub property rented. An owner may own several properties.
- Properties are subject to damage and the agency records all instance of damage to its properties - property, date, type of damage and repair cost are recorded. Repair costs are charged directly to tenants
- Normal property maintenance is also noted - property, date, type of maintenance and cost are recorded. Maintenance costs are charged to the property owner.
- Tenants pay accounts to the Agency - these consist of weekly rental payments, bond payments (for new properties) and damage bills. The date of payment, rental, property, type of account (Rental, Bond, Damage) and amount are recorded.


4.2 SQL Developer Data Modeller Issues

This document lists some of the issues you may experience when using SQL Developer and current work around's

Relational Model Disappears

When a model is saved SQL Developer Data Modeller sometimes fails to save your relational model fully. The Relational diagram is still present in your model but does not show within your project when you re open it.

Before proceeding please ensure your model is closed and you have exited from SQL Developer. To correct this problem, open your model *folder* and navigate to the rel folder (this is where your relational model is saved):



The model shown above is complete, it does not have the missing relational diagram problem. The rel folder should contain a folder and an xml document of the same name. If there is a folder inside rel (it will have a different name, they are all unique) but no xml file of the same name then this is a save error.

Your relational model has not gone, the problem is that the xml document was not saved correctly, follow the steps below to correct this problem:

Make sure your model is closed

1. Copy the [linked XML file](https://goo.gl/Y5Zkzo) (<https://goo.gl/Y5Zkzo>) into the rel folder
2. Rename the file from "CHANGE_TO_FOLDER_NAME.xml" to have the same name as the folder in your rel folder, with the .xml extension added eg. as above 1D583D53-B2C458D2350C.xml (yours will be different)
3. Open the file in a *text editor* and change the id="CHANGE_TO_FOLDER_NAME" on line 2 to be the same as your folder name

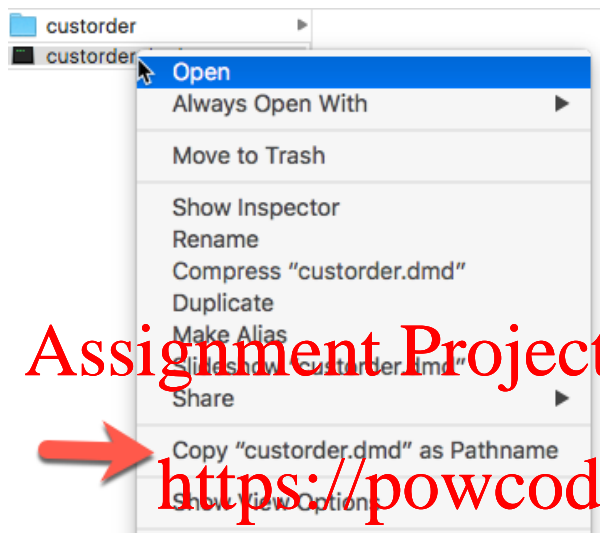
When you now open your model the relational model should be back.

Having problems navigating into a folder

Sometimes the java based file manager of SQL Developer Data Modeler has problems navigating through and selecting your folders.

If this occurs, or as a standard alternative, you can type/obtain in the full path name that you wish to open or navigate to. For example, if you have a Customer-Orders project you wish to open use the following steps:

In your file manager navigate to where your .dmd file for the project is and obtain it's full path name. Under Mac OS, right click the dmd file and right click and while holding this, press the "option" key, this

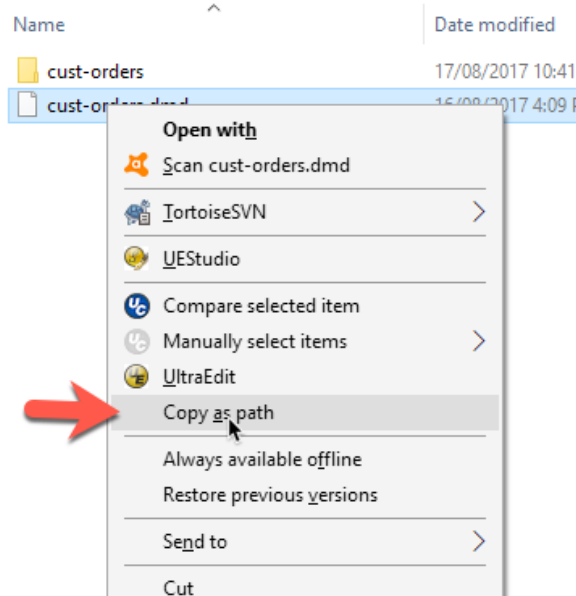


Assignment Project Exam Help

<https://powcoder.com>


will result in

under MS Windows, hold the shift and while holding it, right click on the file or folder, "Copy as path" will



appear

This path can then be pasted into the SQL Developer dialog box - to cause it to go to a particular folder or open a particular project. For example, in MS Windows:



Note that in MS Windows you must first remove the opening and closing quotes (") from the path name.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

5

Week 5: Normalisation

Normalisation Tutorial Activities

Introduction

Normalisation process starts with identifying attributes and the possible existence of the repeating group. The starting set of attributes is called UNF (un-normalised form). During analysis, it is possible to have different sets of un-normalised form (UNF). This initial form will depend on what you perceive being represented in the table or relation. For example, consider the following table. You may consider the table is about PROJECT, others may see the table is data about PROJECT MANAGER. Both are correct perceptions but it will lead to a different UNF in the normalisation process.

	PROJECT_CODE	PROJECT_MANAGER	MANAGER_PHONE	MANAGER_ADDRESS	PROJECT_BID_PRICE
▶	21-5Z	Holly B. Parker	904-338-3416	3334 Lee Rd., Gainesville, FL 37123	\$16,833,460.00
	25-2D	Jane D. Grant	615-898-9909	218 Clark Blvd., Nashville, TN 36362	\$12,500,000.00
	25-5A	George F. Dorts	615-227-1245	124 River Dr., Franklin, TN 29185	\$32,512,420.00
	25-9T	Holly B. Parker	904-338-3416	3334 Lee Rd., Gainesville, FL 37123	\$21,563,234.00
	27-4Q	George F. Dorts	615-227-1245	124 River Dr., Franklin, TN 29185	\$10,314,545.00
	29-2D	Holly B. Parker	904-338-3416	3334 Lee Rd., Gainesville, FL 37123	\$25,559,999.00
	31-7P	William K. Moor	904-445-2719	216 Morton Rd., Stetson, FL 30155	\$56,850,000.00

A. UNF when perceiving table to represent PROJECT

PROJECT(project_code, project_manager, manager_phone, manager_address, project_bid_price)

B. UNF when perceiving table to represent PROJECT MANAGER

MANAGER(project_manager, manager_phone, manager_address (project_code, project_bid_price))

When you consider the table represents PROJECT MANAGER, you see that we have a repeating group because a PROJECT MANAGER can supervise more than one PROJECT. Notice the curly brackets enclosing project_code and project_bid_price attributes.

The PROJECT table/relation, on the other hand, does not have a repeating group because a project has only one project manager.

To complete the normalisation, you choose one of the UNFs as a starting point and follow the normalisation step to at least 3NF. The example steps to complete the normalisation provided below.

Normalisation process:

Legend: Primary Key is underlined, Foreign Key is in italic.

The identification of Foreign Key is not part of the formal steps of normalisation, however, it will give you a mechanism to check whether you have split the attributes of a relation into two separate relations correctly. When you split a relation into two, there should be at a common attribute (or attributes) exists in both relations. In one relation, it will be a Primary Key (PK). In the other relation, the attribute will be a Foreign Key (FK).

Approach A.

UNF

PROJECT(project_code, project_manager, manager_phone, manager_address, project_bid_price)

1NF

PROJECT(project_code, project_manager, manager_phone, manager_address, project_bid_price)

2NF

There is no partial dependency, the 2NF is the same with the 1NF.

3NF

PROJECT(project_code, project_bid_price, *project_manager*)

MANAGER(project_manager, manager_phone, manager_address)

Approach B

UNF

MANAGER(project_manager, manager_phone, manager_address (project_code, project_bid_price))

1 NF

MANAGER(project_manager, manager_phone, manager_address)

PROJECT(project_code, project_bid_price, *project_manager*)

2 NF

There is no partial dependency, the 2 NF is the same to the 1NF.

3 NF

There is no transitive dependency, the 2NF is the same to the 2NF.

Tutorial Questions.

A. This exercise is adapted from Connolly and Begg, Exercise 14.15 (p441)

staffNo	dentistName	patNo	patName	appointment		surgeryNo
				date	time	
S1011	Tony Smith	P100	Gillian White	12-Sep-08	10:00	S15
S1011	Tony Smith	P105	Jill Bell	12-Sep-08	12:00	S15
S1024	Helen Pearson	P108	Ian MacKay	12-Sep-08	10:00	S10
S1024	Helen Pearson	P108	Ian MacKay	14-Sep-08	14:00	S10
S1032	Robin Plevin	P105	Jill Bell	14-Sep-08	16:30	S15
S1032	Robin Plevin	P110	John Walker	15-Sep-08	18:00	S13

<https://powcoder.com>

- Describe possible insertion, deletion and update anomalies for the DENTIST relation.
- Write the UNF (un-normalised Form) for each of the possible analysis below for the data depicted on the above table.
 - The table contains data about APPOINTMENTS. An appointment is made for a dentist and a patient at an agreed appointment time.
 - The table contains data about DENTISTS. A dentist has a number of appointments.
 - The table contains data about DENTISTS. A dentist sees a number of patients. For each patient it may involve several appointments.
- Perform the normalisation process up to the 3NF for the UNF you have created based on statement 1. After you produce the 1NF (the first normal form), draw the dependency diagram of the relation(s).
- Perform the normalisation process up to the 3NF for the UNF you have created based on statement 2. After you produce the 1NF (the first normal form), draw the dependency diagram of the relation(s).

B. Normalise the forms to the 3NF.

UNITS CURRENTLY APPROVED

REPORT

DATE:

9/07/2015

Unit Number	Unit Name	Unit Description	Unit Value
FIT9131	Programming Foundations	Introduction to programming	6
FIT9132	Introduction to Databases	Database Fundamentals	6
FIT9134	Computer Architecture and Operating Systems	Fundamentals of computer systems and the computing environment	6

FIT9135 Data Communications Fundamentals of data and computer communications 6

* Unit values may be either 3, 6 or 12 points

LECTURER DETAILS

REPORT DATE: 29/07/2015

LECTURER'S NUMBER: 10234

LECTURER'S NAME: GUISEPPE BLOGGS

LECTURER'S OFFICE No.: 169

LECTURER'S PHONE No.: 99037111

UNIT ADVISER FOR:

UNIT NUMBER UNIT NAME

FIT9131 Programming Foundations

FIT9134 Computer Architecture and Operating Systems

* A given unit may have several advisers

* Some lecturers share offices, although each have their own phone

STUDENT DETAILS

REPORT DATE: 29/07/2015

STUDENT No.: 12345678

STUDENT NAME: Poindexter Jones

STUDENT ADDRESS: 23 Wide Road, Caulfield, 3162

COURSE ENROLLED: MIT

MODE OF STUDY: On-Campus

MENTOR NUMBER: 10234

MENTOR NAME: Giuseppe Bloggs

ACADEMIC RECORD:

UNIT NUMBER UNIT NAME YEAR / SEMESTER GRADE

FIT9131 Programming Foundations 2014/2 N

FIT9131 Programming Foundations 2015/1 S

FIT9132 Introduction to Databases 2015/1 D

* Grade may have the value N, P, C, D or HD

* Mode of Study must be On-campus (O) or Distance Education (D)

In order to add a student, the lecturer who advises that student must already exist in the database. No lecturer may be deleted who advises any students which are currently in the database. If the lecturer number of a lecturer is changed, then the number would be changed for each student advised by that lecturer.

Note:

- To simplify the normalisation process that involves multiple forms, you should perform the normalisation one form at a time until all relations are in 3NF. Once you have done this process for all forms, consolidate the relations from the different forms by:
 - group together all relations with the same primary key, i.e representing the same entity.
 - choose a single name for synonyms. For example, **mentor** is the same as **lecturer**.

6

Week 6: SQL Data Definition Language (DDL)

When creating schema files, you should always also create a drop file or add the drop commands to the top of your schema file. You should drop the tables using the:

drop table tablename purge;


syntax.

The drop table statements should list tables in the reverse order of your create table order so that FK relationships will be able to be removed successfully. Should a syntax error occur while testing your schema, you simply need to run the drop commands to remove any tables which may have been created.

An excellent summary of the Oracle data types and version restrictions is available from:

<https://www.techonthenet.com/oracle/datatypes.php>

For this unit, we make use of CHAR, VARCHAR2 (or VARCHAR), NUMBER (or NUMERIC), and DATE



The data model above represents figure 3.3 from Coronel & Morris. There are two different ways of coding this model as a set of create table statements.

Using table constraints

SQL constraints are classified as column or table constraints; depending on which item they are attached to:

```
create table agent
(
  agent_code      number (3) constraint agent_pk primary key,
  agent_areacode  number (3) not null ,
  agent_phone     char (8) not null ,
  agent_lname     varchar2 (50) not null ,
  agent_ytd_sls   number (8,2) not null
);
```

This is a declaration of the primary key as a column constraint

```

create table agent
(
    agent_code      number (3) not null ,
    agent_areacode  number (3) not null ,
    agent_phone     char (8) not null ,
    agent_lname     varchar2 (50) not null ,
    agent_ytd_sls   number (8,2) not null,
    constraint agent_pk primary key ( agent_code )
) ;

```

Here the primary key has been declared as a table constraint, at the end of the table after all column declarations have been completed. In some circumstances, for example, a composite primary key you must use a table constraint since a column constraint refers only to a single column.

The create table statements for the two tables in fig 3-3 would be:

```

create table agent
(
    agent_code      number (3) not null ,
    agent_areacode  number (3) not null ,
    agent_phone     char (8) not null ,
    agent_lname     varchar2 (50) not null ,
    agent_ytd_sls   number (8,2) not null,
    constraint agent_pk primary key ( agent_code )
) ;

```

```

create table customer
(
    cus_code        number (5) not null ,
    cus_lname       varchar2 (50) not null ,
    cus_fname       varchar2 (50) not null ,
    cus_initial     char (1) ,
    cus_renew_date  date not null ,
    agent_code      number (3),
    constraint customer_pk primary key ( cus_code ),
    constraint customer_agent_fk foreign key ( agent_code)
        references agent ( agent_code ) on delete set null
) ;

```

In some circumstances, this approach cannot be used. Can you see what the issue is with trying to create the two tables depicted below?



In such a situation an alternative approach to declaring constraints needs to be adopted.

Using ALTER table commands

In this approach, the tables are declared without constraints and then the constraints are applied via the ALTER TABLE command (see section 7.5 of Coronel & Morris).

```
create table agent
(
    agent_code      number (3) not null ,
    agent_areacode  number (3) not null ,
    agent_phone     char (8) not null ,
    agent_lname     varchar2 (50) not null ,
    agent_ytd_sls   number (8,2) not null
);

alter table agent add constraint agent_pk primary key
    ( agent_code );
```

```
create table customer
(
    cus_code        number (5) not null ,
    cus_lname       varchar2 (50) not null ,
    cus_fname       varchar2 (50) not null ,
    cus_initial     char (1) ,
    cus_renew_date  date not null ,
    agent_code      number (3)
);

alter table customer add constraint customer_pk primary key
    ( cus_code );

alter table customer add constraint customer_agent_fk foreign key
    ( agent_code ) references agent ( agent_code )
    on delete set null;
```

Assignment Project Exam Help
<https://powcoder.com>
Add WeChat powcoder

After creating the tables we need to insert the data, for AGENT the insert will have the form:

```
insert into agent values (501,713,'228-1249','Alby',132735.75);
```


for customer:

```
insert into customer values(10010,'Ramas','Alfred','A',
    '05-Apr-2014',501);
```

It is important to note that for the insert into customer we are using the default Oracle date format of dd-mon-yyyy - in the near future we will correct this and allow any date format via the oracle function *to_date*.

Lab Tasks

Using the model from the DDL lecture for student, unit and enrolment:



Creating tables from scratch.

- Code a schema file to create these three tables, noting the following extra constraints:
 - stu_nbr > 10000000
 - unit_name is unique in the UNIT table
 - enrol_semester can only contain the value of 1 or 2 or 3.
- In implementing these constraints you will need to make use of CHECK clauses (see Coronel & Morris section 7.2.6).
- Ensure your script file has appropriate comments in the header, includes the required drop commands and includes echo on and echo off commands.
- Run your script and create the three required tables.
- Save the output from this run.

As an alternative to using echo on/off and having to save the output, a simpler approach is through the use of the inbuilt Oracle SPOOL command.

To use SPOOL, place as the top line in your schema file:

```
spool ./myoutput.txt
```

and as the last line in your script file

```
spool off
```

This will produce a file, in the same folder that your script is saved in, called myoutput.txt which contains the full run of your SQL script.

6.1 INSERTing data into the database

1. Basic INSERT statement.

In this exercise, you will enter the data into the database using INSERT statements with the following assumptions:

- the database currently does not have any existing data.
- the primary key is not generated automatically by the DBMS.

TASKS

Insert the following data into the tables specified using the SQL INSERT statement:

STUDENT

stu_nbr	stu_lname	stu_fname	stu_dob
11111111	Boggs	Fred	01-Jan-1990
11111112	Nice	Nick	10-Oct-1994
11111113	Wheat	Wendy	05-May-1990
11111114	Steen	Clara	25-Dec-1991

UNIT

unit_code	unit_name
FIT9999	FIT Last Unit
FIT5132	Introduction to Databases
FIT5016	Project
FIT5111	Student's Life

ENROLMENT

stu_nbr	unit_code	enrol_year	enrol_semester	enrol_mark	enrol_grade
11111111	FIT5132	2013	1	35	N
11111111	FIT5016	2013	1	61	C
11111111	FIT5132	2013	2	42	N
11111111	FIT5111	2013	2	76	D
11111111	FIT5132	2014	2		
11111112	FIT5132	2013	2	83	HD
11111112	FIT5111	2013	2	79	D
11111113	FIT5132	2014	2		
11111113	FIT5111	2014	2		
11111114	FIT5111	2014	2		

- Ensure you make use of COMMIT to make your changes permanent.
- Check that your data has inserted correctly by using the SQL command `SELECT * FROM tablename` and by using the SQL GUI (select the table in the right-hand list and then select the Data tab).

2. Using SEQUENCES in an INSERT statement.

In the previous exercises, you have entered the primary key value manually in the INSERT statements. In the case where a SEQUENCE is available, you should use the sequence mechanism to generate the value of the primary key.

TASKS

Create a sequence for the STUDENT table called STUDENT_SEQ

- Create a sequence for the STUDENT table called STUDENT_SEQ that starts at 11111115 and increases by 1.
- Check that the sequence exists in two ways (using SQL and browsing your SQL Developer connection objects).

Add a new student (MICKEY MOUSE)

- Use the student sequence - pick any STU_DOB you wish.
- Check that your insert worked.
- Add an enrolment for this student to the unit FIT5132 in semester 2 2016.

Assignment Project Exam Help

3. Advanced INSERT.

We have learned how to add data into the database in the previous exercises through the use of INSERT statements. In those exercises, the INSERT statements were created as a single script assuming that data is all added at the same time, such as at the beginning when the tables are created. On some occasions, new data is added after some data already exists in the database. In this situation, it is a good idea to use a combination of INSERT and SELECT statements.

<https://powcoder.com>

Add WeChat powcoder


A SELECT statement is an SQL statement that we use to retrieve data from a database. An example of a SELECT statement would be:

```
SELECT vendor_id
FROM vendor
WHERE vendor_name = 'Seagate';
```

The above SQL statement consists of three SQL clauses SELECT, FROM and WHERE. The SELECT clause is used to declare which column(s) are to be displayed in the output. The FROM clause is used to declare from which table the data needs to be retrieved. The WHERE clause is used to declare which rows are to be retrieved. In the above SQL select, any row that has the vendor_name equal to 'Seagate' will be retrieved. The SQL SELECT statement will be covered in more detail in the future module, retrieving data from the database.

For our exercise on using the advanced INSERT statement, consider the following model depicting VENDOR and PRODUCT.

Assume we want to add vendors and the products they supply into a set of tables represented by:



A suitable schema would be:

```
DROP TABLE PRODUCT PURGE;
DROP TABLE VENDOR PURGE;
DROP SEQUENCE PRODUCT_prod_no_SEQ;
DROP SEQUENCE VENDOR_vendor_id_SEQ;
```

```
CREATE TABLE PRODUCT
(
    prod_no          NUMBER (4) NOT NULL ,
    prod_name        VARCHAR2 (50) NOT NULL ,
    prod_price       NUMBER (6,2) NOT NULL ,
    prod_stock       NUMBER (3) NOT NULL ,
    VENDOR_vendor_id NUMBER (3) NOT NULL
);
```

```
ALTER TABLE PRODUCT ADD CONSTRAINT PRODUCT_PK PRIMARY KEY ( prod_no );
```

```
CREATE TABLE VENDOR
(
    vendor_id        NUMBER (3) NOT NULL ,
    vendor_name      VARCHAR2 (50) NOT NULL ,
    vendor_phone     CHAR (10) NOT NULL
);
```

```
ALTER TABLE VENDOR ADD CONSTRAINT VENDOR_PK PRIMARY KEY ( vendor_id );
```

```
ALTER TABLE VENDOR ADD CONSTRAINT VENDOR_UN UNIQUE ( vendor_name );
```

```
ALTER TABLE PRODUCT ADD CONSTRAINT PRODUCT_VENDOR_FK FOREIGN KEY (
VENDOR_vendor_id ) REFERENCES VENDOR ( vendor_id ) ON
DELETE CASCADE ;
```

```
CREATE SEQUENCE PRODUCT_prod_no_SEQ START WITH 1 INCREMENT BY 1;
```

```
CREATE SEQUENCE VENDOR_vendor_id_SEQ START WITH 1 INCREMENT BY 1;
```

There are two ways in which we can perform the INSERT.

1. Use the nextval and currval of the sequences.

```
-- Add Vendor 1 and the products they supply
insert into vendor values (VENDOR_vendor_id_SEQ.nextval,
    'Western Digital', '1234567890');
insert into product values (PRODUCT_prod_no_SEQ.nextval,
    '2TB My Cloud Drive',195,5,VENDOR_vendor_id_SEQ.currval);
```

```

insert into product values (PRODUCT_prod_no_SEQ.nextval,
    '1TB Portable Hard Drive',76,4,VENDOR_vendor_id_SEQ.currval);
insert into product values (PRODUCT_prod_no_SEQ.nextval,
    'Live Media Player',119,2,VENDOR_vendor_id_SEQ.currval);
commit;
-- Add Vendor 2 and the products they supply
insert into vendor values (VENDOR_vendor_id_SEQ.nextval,'Seagate',
    '2468101234');
insert into product values (PRODUCT_prod_no_SEQ.nextval,
    '2TB Desktop Drive',94,12,VENDOR_vendor_id_SEQ.currval);
insert into product values (PRODUCT_prod_no_SEQ.nextval,
    '4TB 4 Bay NAS',76,4,VENDOR_vendor_id_SEQ.currval);
insert into product values (PRODUCT_prod_no_SEQ.nextval,
    '2TB Central Personal Storage' ,169,5,
    VENDOR_vendor_id_SEQ.currval);
commit;

```

2. Use the nextval in combination with the SELECT statement.

- Add a new product for a vendor at a subsequent time (vendor names will be unique - note the U in the model above and the vendor_un constraint in the schema)

```

insert into product values (PRODUCT_prod_no_SEQ.nextval,
    'GoFlex Thunderbolt Adaptor',134,12,
    (select vendor_id from vendor where vendor_name = 'Seagate'));

```

In subsequent weeks you will see that the same concept can be used with other data manipulation statements such as UPDATE and DELETE.

TASKS

Add WeChat powcoder

- A new student has started a course and needs to enrol into "Introduction to databases". Enter the new student's details and his/her enrolment to the database using the nextval in combination with a SELECT statement. You can make up details of the new student and when they will attempt "Introduction to databases".
 - You must not do a manual lookup to find the unit code of the "Introduction to databases".

4. Creating a table and inserting data as a single SQL statement.

A table can also be created based on an existing table, and immediately populated with contents by using a SELECT statement within the CREATE TABLE statement.

For example, to create a table called FIT5132_STUDENT which contains the enrolment details of all students who have been or are currently enrolled in FIT5132, we would use:

```

create table FIT5132_STUDENT
as select *
from enrolment
where unit_code = 'FIT5132';

```

Here, we use the SELECT statement to retrieve all columns (the wildcard "*" represents all columns) from

the table enrolment, but only those rows with a value of the unit_code equal to FIT5132.

TASKS

- Create a table called FIT5111_STUDENT. The table should contain all enrolments for the unit FIT5111.
- Check the table exists.
- List the contents of the table.

5. Changing a table's structure.

TASKS

- Add a new column to the UNIT table which will represent credit points for the unit (hint use the ALTER command). The default value for the new column should be 6 points.
- Insert a new unit after you have added the new column. You can make up the details of the new unit.
- Check that the new insert has worked correctly.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

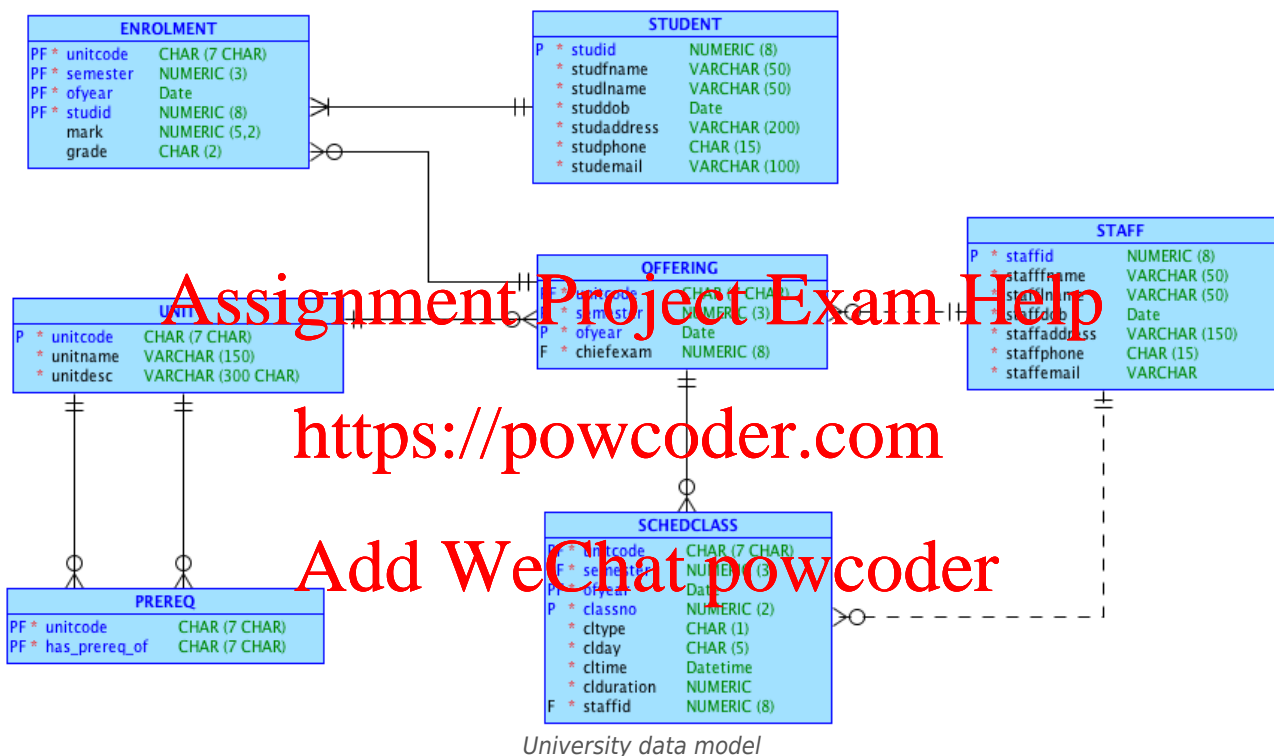
7

Week 7: SQL Part I - SQL Basic

The following exercises will allow you to be familiar with writing basic SQL statements. Two main concepts that you need to learn from these exercises are:

- Using the WHERE clause to retrieve selected rows from a table.
- Using the WHERE clause to join two or more tables.

Use the UNIVERSITY database to complete the exercises. Figure 1.0 depicts the data model for the UNIVERSITY database.



For students using the Monash Oracle database, to complete these exercises, you do not need to run the supplied DDL scripts in your account. A UNIVERSITY database has been created under the user "UNI". To use the database you need to include the prefix "UNI" to the table names that you use in an SQL statement. For example, if you want to retrieve data from UNIT table you need to write:

```
SELECT unitcode, unitname
FROM uni.unit;
```

instead of

```
SELECT unitcode, unitname
FROM unit;
```

This week we make use of Oracle dates - to use these correctly you should note the following:

- The Oracle date data type contains both date and time, however you can choose to use just a date, just a time, both or parts of a date depending on the format strings used
- `to_date`: converts from a string to a date according to a format string
- `to_char`: converts from a date to a string according to a format string

The Oracle documentation links are:

- [Format models](http://goo.gl/6IFTqP) (<http://goo.gl/6IFTqP>)
- [to_date](http://goo.gl/hhrCAo) (<http://goo.gl/hhrCAo>)
- [to_char](http://goo.gl/7yvmwp) (<http://goo.gl/7yvmwp>)

Part A. Retrieving data from a single table

1. List all students and their details.
2. List all units and their details.
3. List all students who have the surname 'Smith'.
4. List the student's details for those students who have surname starts with the letter "S". In the display, rename the columns `studfname` and `studlname` to `firstname` and `lastname`.
5. List the student's surname, firstname and address for those students who have surname starts with the letter "S" and firstname contains the letter "i".
6. List the unit code and semester of all units that are offered in the year 2014.
To complete this question you need to use the Oracle function `to_char` to convert the data type for the year component of the offering date into text. For example, `to_char(ofyear,'yyy')` - here we are only using the year part of the date.
7. List the unit code of all units that are offered in semester 1 of 2014.
8. Assuming that a unit code is created based on the following rules:
 - The first three letters represent faculty abbreviation, eg FIT for the Faculty of Information Technology.
 - The first digit of the number following the letter represents the year level. List the unit details of all first year units in the Faculty of Information Technology.
9. List the unit code and semester of all units that were offered in either semester 1 or summer of 2013. Note: summer semester is recorded as semester 3.
10. List the student number, mark, unit code and semester for those students who have passed any unit in semester 1 of 2013.

Part B. Retrieving data from multiple tables.

Note: remember to use the foreign key and the primary key when joining two or more tables.

1. List the name of all students who have marks in the range of 60 to 70.
2. List all the unit codes, semester and name of the chief examiner for all the units that are offered in 2014.
3. List the name (firstname and surname), unit names, the year and semester of enrolment of all units taken so far.
4. List all the unit codes and the unit names and their year and semester offerings. To display the date correctly in Oracle, you need to use `to_char` function. For example, `to_char(ofyear,'YYYY')`.
5. List the unit code, semester, class type (lecture or tutorial), day and time for all units taught by Albus Dumbledore in 2013. Sort the list according to the unit code.
6. Create a study statement for Mary Smith. A study statement contains unit code, unit name, semester and year study was attempted, the mark and grade.
7. List the unit code, unit name and the unit code and unit name of the pre-requisite units of all units in the database.

8. List the unit code and unit name of the pre-requisite units of 'Advanced Data Management' unit.
 9. Find all students (list their id, firstname and surname) who have a failed unit in the year 2013.
 10. List the student name, unit code, semester and year for those students who do not have marks recorded.
-

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

8

Week 8: Maintaining data

Using DELETE and UPDATE

1. UPDATE

It is common for data to change value across time. In a database, we use the SQL UPDATE statement to change the value of a cell or cells in a table.

The UPDATE statement consist of three main components:

- The name of the table where the data will be updated.
- The row or the set of rows where the value will be updated.
- The new value to replace the old value.

An example of an UPDATE statement for data in the database we have created by following the exercises in the module [Tutorial 6 SQL Data Definition Language DDL](#) is as follow:

```
UPDATE enrolment
SET enrol_mark = 60
WHERE stud_nbr = 11111111 AND
unit_code = 'FIT5132' AND
enrol_semester = '2' AND
enrol_year = 2014;
```

Assignment Project Exam Help

<https://powcoder.com>

TASKS

Add WeChat powcoder

1. Update the unit name of FIT9999 from 'FIT Last Unit' to 'place holder unit'.
2. Enter the mark and grade for the student with the student number of 11111113 for the unit code FIT5132 that the student enrolled in semester 2 of 2014. The mark is 75 and the grade is D.
3. The university introduced a new grade classification. The new classification are:
 1. 45 - 54 is P1.
 2. 55 - 64 is P2.
 3. 65 - 74 is C.
 4. 75 - 84 is D.
 5. 85 - 100 is HD.

Change the database to reflect the new grade classification.

2. DELETE

The DELETE statement is used to remove data from the database. It is important to consider the referential integrity issues when writing a DELETE statement. In Oracle, a table can be created with FOREIGN KEY constraints with a reference_clause ON DELETE action. The action can be set to CASCADE and SET NULL. When the reference_clause is not specified, the action will be set to RESTRICT, in other words, the deletion of row in a parent table (table contains the PRIMARY KEY being referred to by a FOREIGN KEY) will not be allowed when there are rows in the child table (the table with the FOREIGN KEY).

TASKS

1. A student with student number 11111114 has taken intermission in semester 2 2014, hence all the enrolment of this student for semester 2 2014 should be removed. Change the database to reflect this situation.
 2. Assume that Wendy Wheat has withdrawn from the university. Remove her details from the database.
 3. Add Wendy Wheat back to the database (use the INSERT statements you have created when completing module [Tutorial 6 SQL Data Definition Language DDL](#)). Change the FOREIGN KEY constraints definition for table STUDENT so it will now include the references_clause ON DELETE CASCADE. Hint: You need to use the ALTER TABLE statement to drop the FOREIGN KEY constraint first and then put it back using ALTER TABLE with ADD CONSTRAINT clause. A brief description of using ALTER to drop a constraint is available [here](https://www.techonthenet.com/oracle/foreign_keys/drop.php) (https://www.techonthenet.com/oracle/foreign_keys/drop.php), the ADD CONSTRAINT was covered in tutorial 6. For more details, you can check the SQL Reference Manual (available from Moodle) for the full syntax and a range of examples. Once you have changed the table, now, perform the deletion of the Wendy Wheat row in the STUDENT table. Examine the ENROLMENT table. What happens to the enrolment records of Wendy Wheat?
-

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

8.1

Transactions Management

Transaction Management

In these exercises, you will examine the issues involved in updating shared data.

You will work in pairs. Suppose one user1 is User1, and the other is User2. *Replace User1 and User2 with your respective usernames when reading the tutorial exercises.*

User1 will create a table called account which will be shared with User2, that is, both users will be allowed to select data from the table and also update data in the table. This table keeps the account balances of customers, where each customer is identified by a unique id.

Q1. User1 should create the account table. The table will have 2 attributes, id and balance. Both attributes will have datatype number. After creating the table, User1 should insert two rows of data so that the table looks as below:

	ID	BALANCE
1	1	100
2	2	200

Assignment Project Exam Help

<https://powcoder.com>

Q2. User1 can now make the account table available to User2 using the following command:

Add WeChat powcoder

```
grant select, update on account to User2;
```

Q3. In order for User2 to access the account table, they would normally have to prefix the account table with the name of the owner, e.g:

```
select * from User1.account;
```

However, we can remove the need to do this if we create a synonym (essentially a system maintained alias for the table) as follows (User2 issues this command):

```
create synonym account for User1.account;
```

Q4. Make sure both users have the autocommit feature turned OFF - i.e. both users should issue the command

```
set autocommit off
```

Q5. Now, try the following (maintain the order of the operations).

- User1 updates the balance of customer 1 from 100 to 110 (without issuing a commit).
- User2 views the contents of the account table (do they see the new value? - if not, why not?)
- User1 issues a commit command.
- User2 views the contents of the account table (do you notice any difference?)

Explain what is happening in the results of the above queries, in the context of atomic transactions.

Q6. Now we will try and see what happens when we try some concurrent updates of the table (keep the order of transactions the same as below)

- User1 updates the balance of customer 2 from 200 to 150 (without issuing a commit).
- User2 tries to update the balance of customer 2 to 100 (what happens?)

Explain what is happening here. What should be done to allow the User2 update to proceed?

Q7. Now try the following:

- User1 updates the balance of customer 2 from 200 to 150 (without issuing a commit).
- User2 tries to update the balance of customer 1 to 125 (what happens?)

How does this differ from the results of the transactions in part 6? What does this tell you about the granularity of locking in Oracle? What must be done in order for the results of both updates to be visible to both users?

Q8. Try and generate a deadlock between the two users (hint: you will need to set up another shared table). Remember that a deadlock occurs when User1 holds a lock on table A and requests a lock on table B, but table B is locked by User2, who is also requesting a lock on table A.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

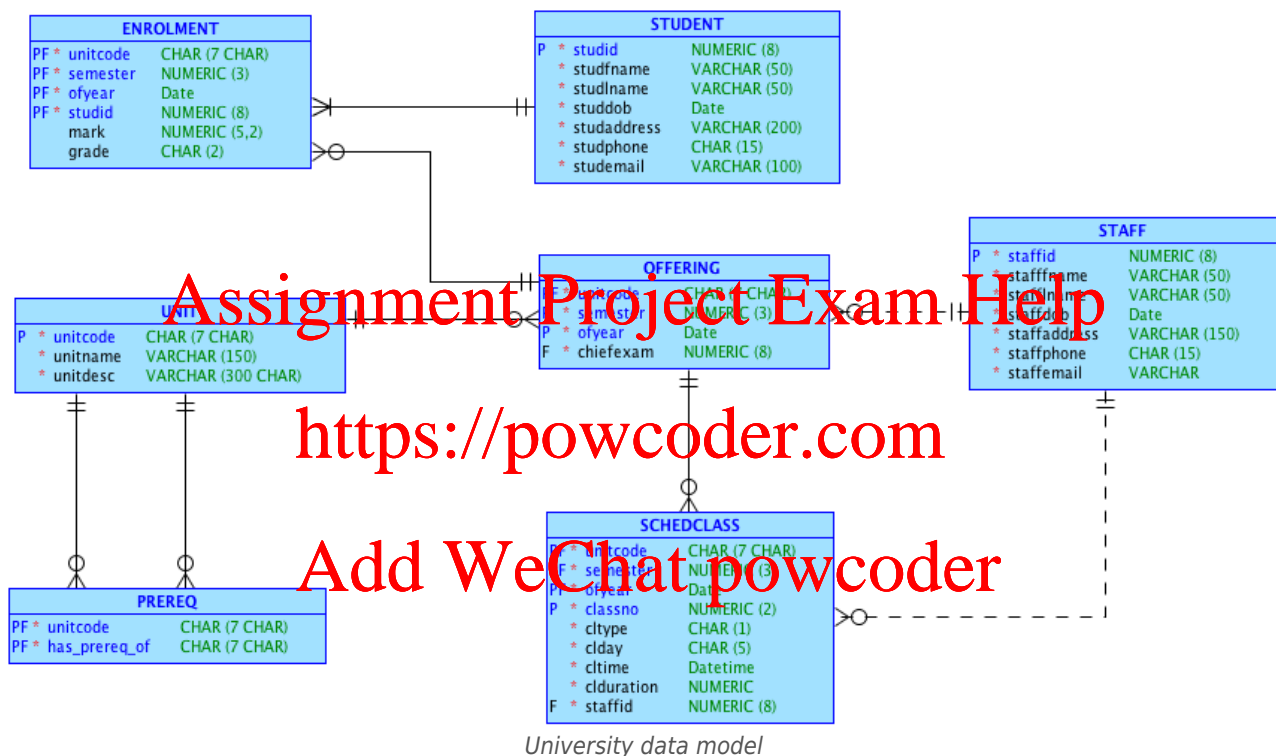
9

Week 9: SQL Part II - SQL Intermediate

The following exercises will allow you to be familiar with:

- All the SQL constructs used in last weeks exercises.
- Aggregate functions such as min, max, avg
- The GROUP BY clause, and
- Subqueries

This week we will continue to use the UNIVERSITY database model:



University data model

Use the Uni database from the previous exercises to complete the following queries.

1. Find the maximum mark for FIT1004 in semester 1, 2013.
2. Find the average mark of FIT1040 in semester 2, 2013.
3. List the average mark for each offering of FIT1040. In the listing, you need to include the year and semester number. Sort the result according to the year.
4. Find the number of students enrolled in the unit FIT1040 in the year 2013, under the following conditions:
 - Repeat students are counted each time
 - Repeat students are only counted once
5. Find the total number of enrolment per semester for each unit in the year 2013. The list should include the unitcode, semester and year. Order the list in increasing order of enrolment numbers.
6. Find the total number of prerequisite units for FIT2077.
7. Find the total number of prerequisite units for each unit. In the list, include the unitcode for which the count is applicable.

8. For each pre-requisite unit, calculate how many times it has been used as prerequisite. Include the name of the prerequisite unit in the listing .
 9. Find the unit with the highest number of enrolments in a given offering in the year 2013.
 10. Who is the oldest student in FIT1004? Display the student full name and the date of birth.
 11. Find all students enrolled in FIT1004 in semester 1, 2013 who have scored more than the average mark of FIT1004 in the same offering? Display the students' name and the mark. Sort the list in the order of the mark from the highest to the lowest.
-

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

10

Week 10: SQL Part III - SQL Advanced

This week we will continue to use the UNIVERSITY database model.

Q1. Find the total number of prerequisite units for each unit. Include in the list the unitcode of units that do not have prerequisite. Hint: use an outer join.

UNITCODE	NO_OF_PREREQ
FIT1004	0
FIT1040	0
FIT2077	2
FIT5131	1
FIT5132	0
FIT5136	1

Q2. Display unitcode and unitname for units that do not have prerequisite.

UNITCODE	UNITNAME
FIT1004	Introduction to Data Management
FIT1040	Programming Fundamentals
FIT5132	Introduction to Databases

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

There are many approaches that you can take in writing an SQL statement to answer this query. You can use the SET OPERATORS, OUTER JOIN and a SUBQUERY. Write SQL statements based on *all* of these approaches.

Q3. List the unit code and the average mark for each offering. Round the average to 2 digits after the decimal points. If the average result is 'null', display the average as 0.00. All values must be shown with two decimal digits:

UNITCODE	YEAR	SEMESTER	AVERAGE
FIT5131	2014	1	0.00
FIT2077	2013	2	64.50
FIT1004	2013	1	65.50
FIT1040	2013	1	67.80
FIT1040	2013	2	71.83
FIT5132	2013	1	74.00
FIT5132	2013	3	74.00
FIT5136	2013	2	75.50
FIT1004	2013	2	75.83
FIT5131	2013	1	77.50

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

11

Week 11: PL/SQL Triggers

PL_SQL Triggers

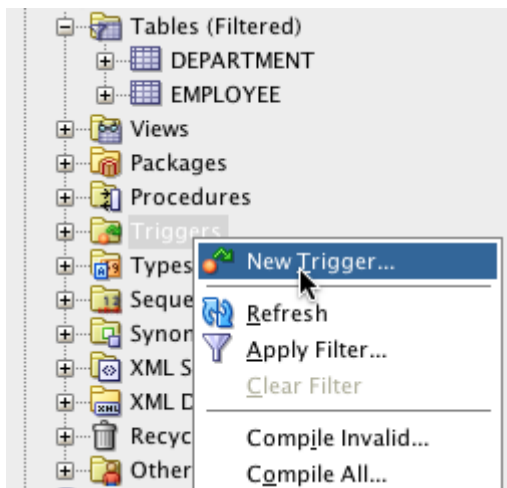
A trigger may be created by typing the code for the trigger directly into the SQL Developer SQL window and finishing the code with a / on a line by itself in column 1, or alternatively via the SQL Developer Trigger item. Using the latter approach allows you to make use of the trigger debug environment in SQL Developer.

Task A: Employee-Department

Given an EMPLOYEE table and a DEPARTMENT table with the structure: DEPARTMENT (dept_no, dept_name dept_empcnt), where dept_empcnt is the count of employees in department. We will create a trigger, which automatically maintains this value:

- insert an employee increases the department count for the department the employee is being added to
- delete an employee decreases the department count for the department the employee is being removed from
- update an employee's department (ie. transfer an employee from their current department to a new department) decreases the current department count and increases the new department count

a. Create a new trigger:



b. Select/Enter the new trigger details:

Schema: LSMI1

Name: MAINTAIN_DEPT_EMP_COUNT

☒ Add New Source In Lowercase

Base Type: TABLE

Base Object Schema: LSMI1

Base Object: EMPLOYEE

Timing: AFTER

Events:

Available Events:	Selected Events:
	DELETE
	INSERT
	UPDATE

Columns:

Available Columns:	Selected Columns:
EMP_NAME	DEPT_NO
EMP_NO	
EMP_SALARY	
EMP_ID	

Referencing Old As: ☐ Statement Level

Referencing New As: When Clause:


Buttons: Help, OK, Cancel

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

c. Select OK and then enter the trigger BODY:



```

1 create or replace trigger maintain_dept_emp_count
2 after delete or insert or update of dept_no on employee
3 for each row
4 begin
5     -- Trigger to automatically maintain the employee count in
6     -- the department table
7
8     -- for inserts, increase employee count in department:
9     IF INSERTING THEN
10        UPDATE DEPARTMENT
11           SET dept_empcnt = dept_empcnt + 1
12           where dept_no = :new.dept_no;
13
14     -- for deletes, decrease employee count in department:
15     ELSE
16        IF DELETING THEN
17           UPDATE DEPARTMENT
18              SET dept_empcnt = dept_empcnt - 1
19              where dept_no = :old.dept_no;
20
21     -- for updates ie moves, modify the employee count in department:
22     ELSE
23        IF UPDATING THEN
24           UPDATE DEPARTMENT
25              SET dept_empcnt = dept_empcnt - 1
26              where dept_no = :old.dept_no;
27           UPDATE DEPARTMENT
28              SET dept_empcnt = dept_empcnt + 1
29              where dept_no = :new.dept_no;
30        END IF;
31     END IF;
32 end;
33


```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

d. finally compile the trigger, correcting any errors which appear:



```

Code | Dependencies | Errors | References | Profiles | Grants | Details
1  c Compile for Debug maintain_dept_emp_cc
2  a Compile update of dept_no or
3  for each row
4  begin
5  Trigger to automatically maintain the

```

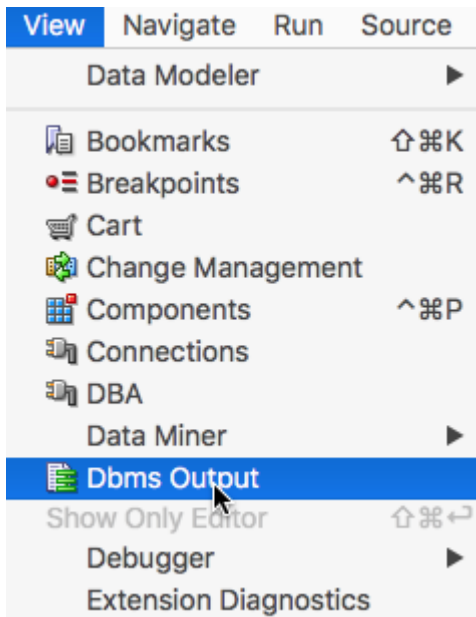
The schema and trigger body code for this example are available in the archive dept-emp.zip available from Moodle.

You can export completed trigger PL/SQL code (right click the trigger and select export) or simply copy and paste the code directly into a SQL script file.


If your output includes a message which will be written to the screen using DBMS_OUTPUT.PUT_LINE you will need to turn on DBMS Output.

To obtain DBMS_OUTPUT in SQL Developer you need to:

(a) Select View - Dbms Output



(b) In the output window which appears, select the green cross and then the connection you wish to receive output from:



Assignment Project Exam Help

Using the archive available from Moodle (dept-emp.zip), create the tables with the schema file, and create the trigger using the above process, copying the trigger's body code from the text file in the archive (copy and paste it into the trigger editor).

<https://powcoder.com>

After you have successfully created the trigger perform the following tasks using SQL:

Add WeChat powcoder

1. Insert two departments, each with an employee count of zero.
2. Now insert an employee and observe what happens with the DEPARTMENT.DEPT_EMPCNT attribute (the count of the employees in a department).
3. Move the employee to the other department (update the employee dept_no), again check the effect on the DEPARTMENT.DEPT_EMPCNT attribute.

Lastly delete the employee, again check the effect on the DEPARTMENT.DEPT_EMPCNT attribute. Drop the employee and department tables.

Task B: Unit-Enrolment-Student

Using the archive student-enrol.zip supplied on Moodle create and load the tables for task B.

1. Create a trigger unit_upd_cascade that updates matching rows in the ENROLMENT table whenever a unit code is updated in the UNIT table. The changes in the enrolment table should be announced via dbms_output.put_line

Check the trigger works using SQL, record your SQL commands and the result. Hint - update a unit code in the unit table and observe what happens in the enrolment table.

2. Create a trigger as discussed in the lecture, to

- Maintain the integrity of the total number of students in a given unit in the UNIT table
- Record a delete operation in the audit_log

Check the trigger works using SQL, record your SQL commands and the result. Hint - enrol and remove an enrolment (delete) for a student.

For students who would like a challenge consider how you might code the trigger to maintain the student's average mark, and what the problems are.

3. Create a trigger called calculate_grade that calculates the student's grade (enrolment.enrol_grade) whenever a mark is updated for an enrolment or a new enrolment is added with a mark. Hint: you can avoid the "mutating table" error here by directly manipulating the new value of enrol_grade in your trigger based on the new enrol_mark value.

- enrol_mark >= 80 is a HD grade
- enrol_mark >= 70 and enrol_mark < 80 is a D grade
- enrol_mark >= 60 and enrol_mark < 70 is a C grade
- enrol_mark >= 50 and enrol_mark < 60 is a P grade
- enrol_mark < 50 is a N grade

Check the trigger works using SQL, record your SQL commands and the result.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder